



# RDBMS & SQL Introduction

Trainer: Mr. Nilesh Ghule

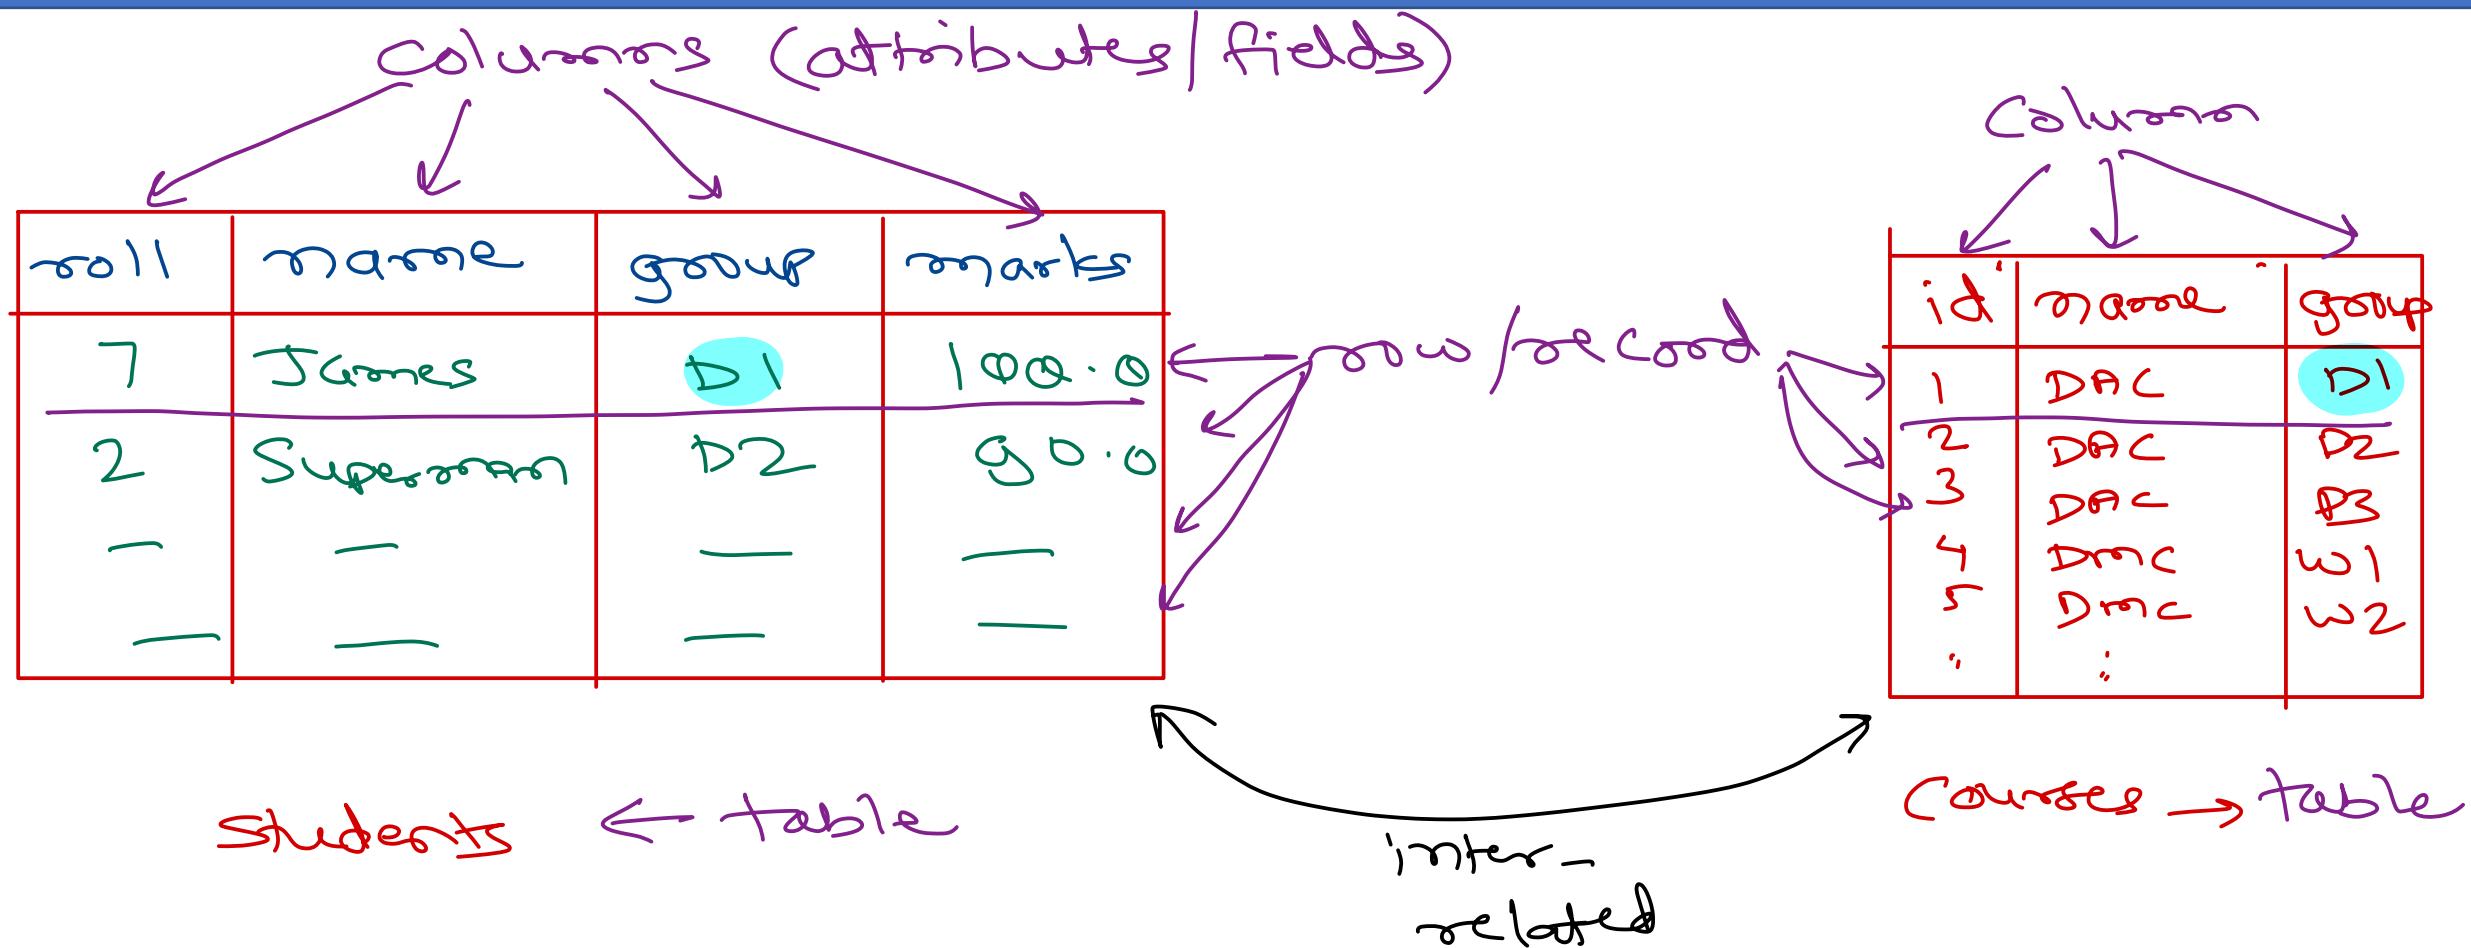


# DBMS

- Any enterprise application need to manage data.
- In early days of software development, programmers store data into files and does operation on it. However data is highly application specific.
- Even today many software manage their data in custom formats e.g. Tally, Address book, etc.
- As data management became more common, DBMS systems were developed to handle the data. This enabled developers to focus on the business logic e.g. FoxPro, DBase, Excel, etc.
- At least CRUD (Create, Retrieve, Update and Delete) operations are supported by all databases.
- Traditional databases are file based, less secure, single-user, non-distributed, manage less amount of data (MB), complicated relation management, file-locking and need number of lines of code to use in applications.



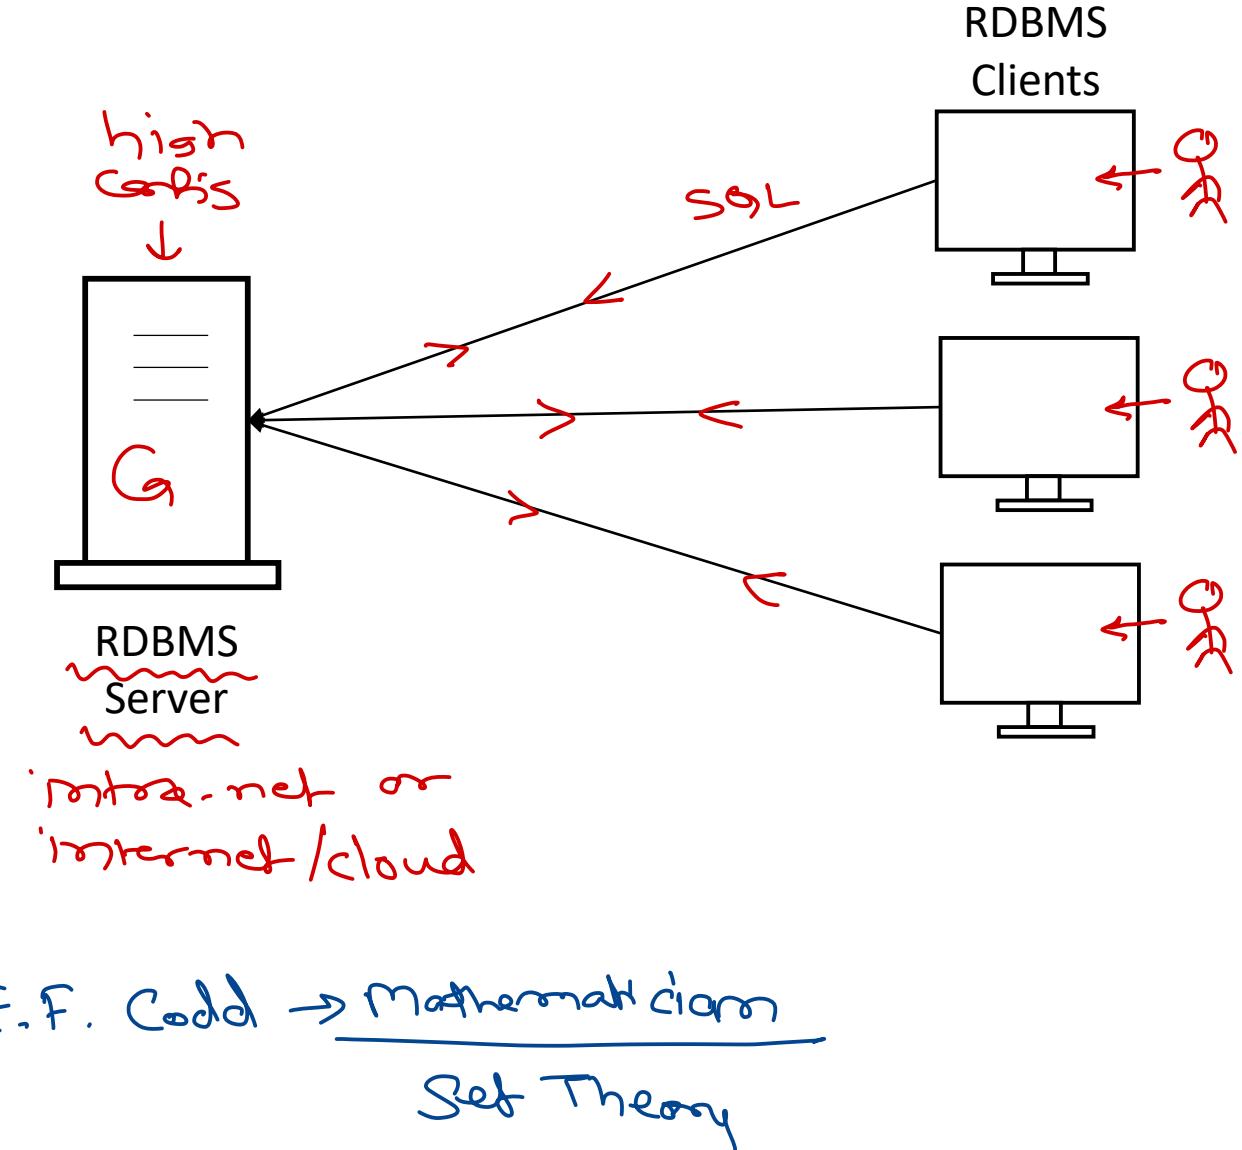
# RDBMS table



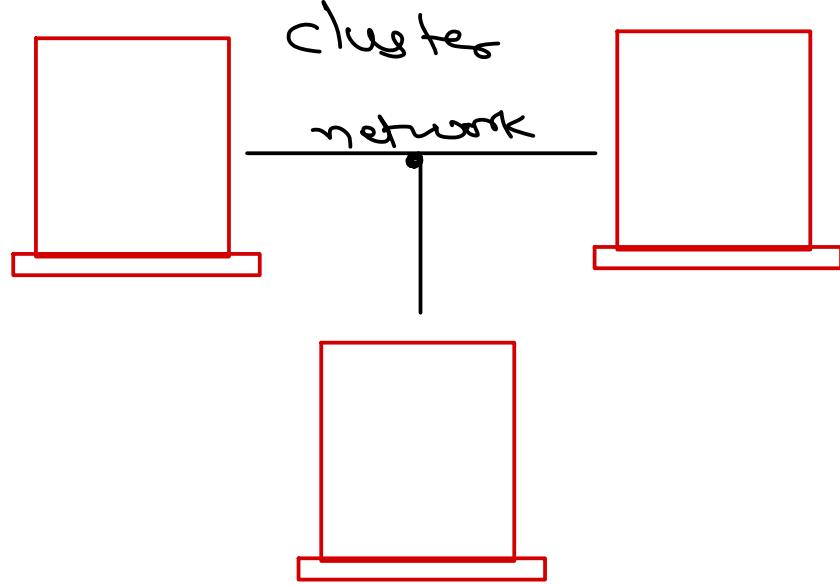
# RDBMS

low  
costs

- RDBMS is relational DBMS.
- It organizes data into Tables, rows and columns. The tables are related to each other. → logically
- RDBMS follow table structure, more secure, multi-user, server-client architecture, server side processing, clustering support, manage huge data (TB), built-in relational capabilities, table-locking or row-locking and can be easily integrated with applications.
- e.g. DB2, Oracle, MS-SQL, MySQL, MS-Access, SQLite, ...
- RDBMS design is based on Codd's rules developed at IBM (in 1970).



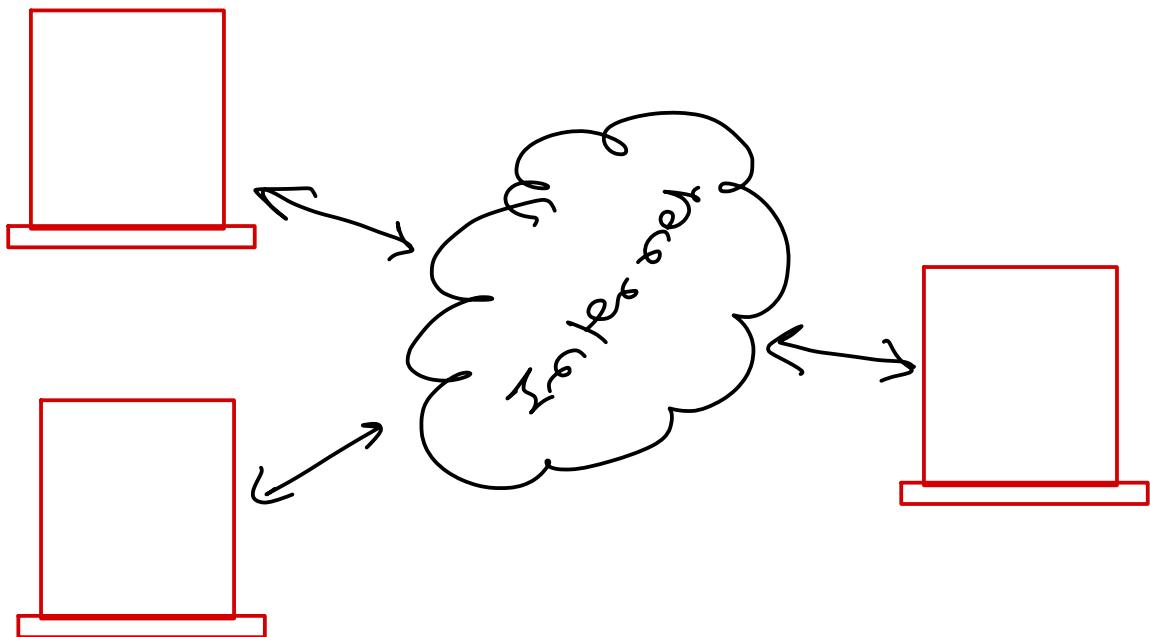
# clustering - distributed systems.



e.g. database cluster  
webserver cluster,

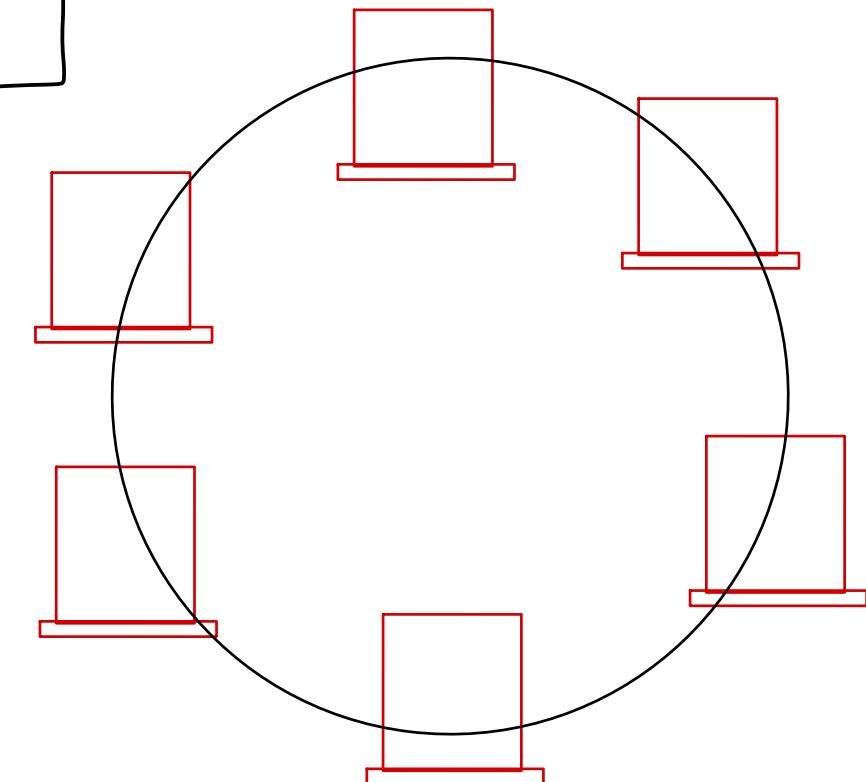
cluster is set of computers  
connected in a network for  
a specific task.

Local net



Internet

Star  
Ring  
Bus

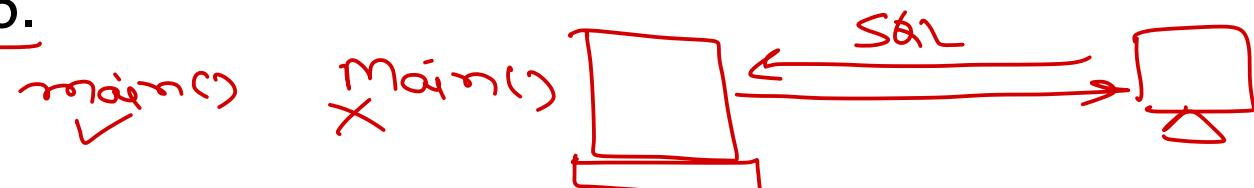


LAN ,MAN ,wAN



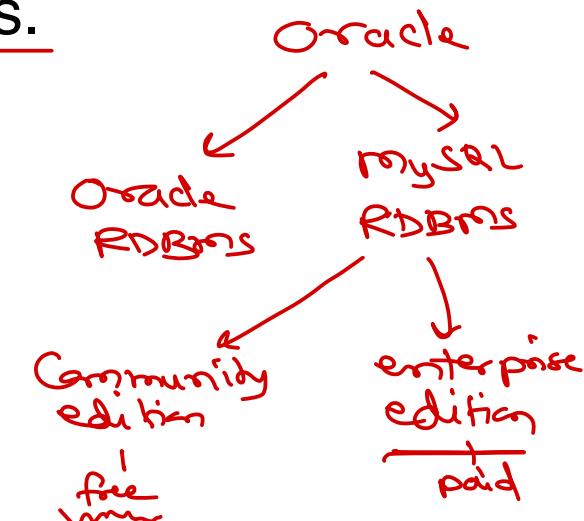
# SQL - Structured Query Language.

- Clients send SQL queries to RDBMS server and operations are performed accordingly.
- Originally it was named as RQBE (Relational Query By Example).
- SQL is ANSI standardised in 1987 and then revised multiple times adding new features. Recent revision in 2016.
- SQL is case insensitive.
- There are five major categories:
  - ✓ DDL: Data Definition Language e.g. CREATE, ALTER, DROP, RENAME.
  - ✓ DML: Data Manipulation Language e.g. INSERT, UPDATE, DELETE.
  - ✓ DQL: Data Query Language e.g. SELECT.
  - ✓ DCL: Data Control Language e.g. CREATE USER, GRANT, REVOKE.
  - ✓ TCL: Transaction Control Language e.g. SAVEPOINT, COMMIT, ROLLBACK.
- Table & column names allows alphabets, digits & few special symbols.
- If name contains special symbols then it should be back-quotes.
  - e.g. Tbl1, `T1#`, `T2\$` etc. Names can be max 30 chars long.



# MySQL - RDBMS Software →(of Oracle)

- Developed by Michael Widenius in 1995. It is named after his daughter name Myia.
- Sun Microsystems acquired MySQL in 2008.
- Oracle acquired Sun Microsystem in 2010.
- MySQL is free and open-source database under GPL. However some enterprise modules are close sourced and available only under commercial version of MySQL.
- MariaDB is completely open-source clone of MySQL.
- MySQL support multiple database storage and processing engines.
- MySQL versions:
  - ✓ < 5.5: MyISAM storage engine
  - ✓ 5.5: InnoDB storage engine
  - ✓ 5.6: SQL Query optimizer improved, memcached style NoSQL
  - ✓ 5.7: Windowing functions, JSON data type added for flexible schema
  - ✓ 8.0: CTE, NoSQL document store.
- MySQL is database of year 2019 (in database engine ranking).



# MySQL installation on Ubuntu/Linux

- terminal> sudo apt-get install mysql-community-server mysql-community-client
- This installs MySQL server (mysqld) and MySQL client (mysql).
- MySQL Server (mysqld)
  - Run as background process.
  - Implemented in C/C++.
  - Process SQL queries and generate results.
  - By default run on port 3306.
  - Controlled via systemctl.
    - terminal> sudo systemctl start|stop|status|enable|disable mysql
- MySQL client (mysql)
  - Command line interface
  - Send SQL queries to server and display its results.
  - terminal> mysql –u root -p
- Additional MySQL clients
  - MySQL workbench
  - PHPMyAdmin





# Thank you!

Nilesh Ghule <[nilesh@sunbeaminfo.com](mailto:nilesh@sunbeaminfo.com)>



# RDBMS & SQL Introduction

Trainer: Mr. Nilesh Ghule



developed in C/C++

mysql -u root -h localhost -p  
user host/server mysql CLI  
← Password

MySQL Server  
mysqld

① accept SQL query

② process query

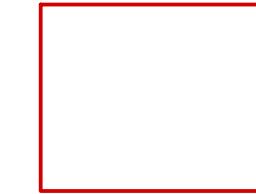
- read data from disk
- process data
- create result

③ return result

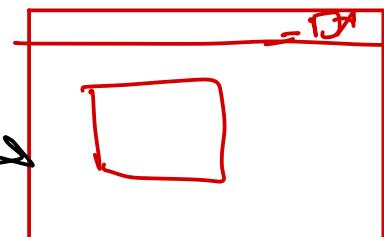
Socket  
IP + Port  
Common endpoint in a network

mysql

mysql shell

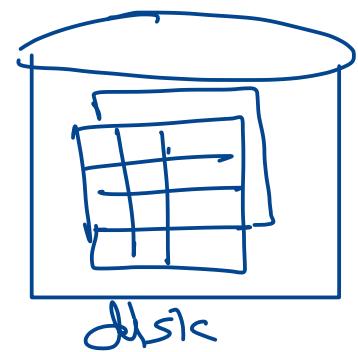


mysql workbench



MySQL daemon

no GUI  
run in background



# MySQL installation on Ubuntu/Linux

- terminal> sudo apt-get install mysql-community-server mysql-community-client
- This installs MySQL server (mysqld) and MySQL client (mysql).
- MySQL Server (mysqld)
  - Run as background process.
  - Implemented in C/C++.
  - Process SQL queries and generate results.
  - By default run on port 3306.
  - Controlled via systemctl.
    - terminal> sudo systemctl start|stop|status|enable|disable mysql
- MySQL client (mysql)
  - Command line interface
  - Send SQL queries to server and display its results.
  - terminal> mysql –u root -p
- Additional MySQL clients
  - MySQL workbench
  - PHPMyAdmin



# Getting started

- root login can be used to perform CRUD as well as admin operations.
- It is recommended to create users for performing non-admin tasks.
  - mysql> CREATE DATABASE db;
  - mysql> SHOW DATABASES;
  - mysql> CREATE USER dbuser@localhost IDENTIFIED BY 'dbpass';
  - mysql> SELECT user, host FROM mysql.user;
  - mysql> GRANT ALL PRIVILEGES ON db.\* TO dbuser@localhost;
  - mysql> FLUSH PRIVILEGES;
  - mysql> EXIT;
- terminal> mysql –u dbuser –pdbpass
  - mysql> SHOW DATABASES;
  - mysql> SELECT USER(), DATABASE();
  - mysql> USE db;
  - mysql> SHOW TABLES;
  - mysql> CREATE TABLE student(id INT, name VARCHAR(20), marks DOUBLE);
  - mysql> INSERT INTO student VALUES(1, 'Abc', 89.5);
  - mysql> SELECT \* FROM student;



# Database logical layout

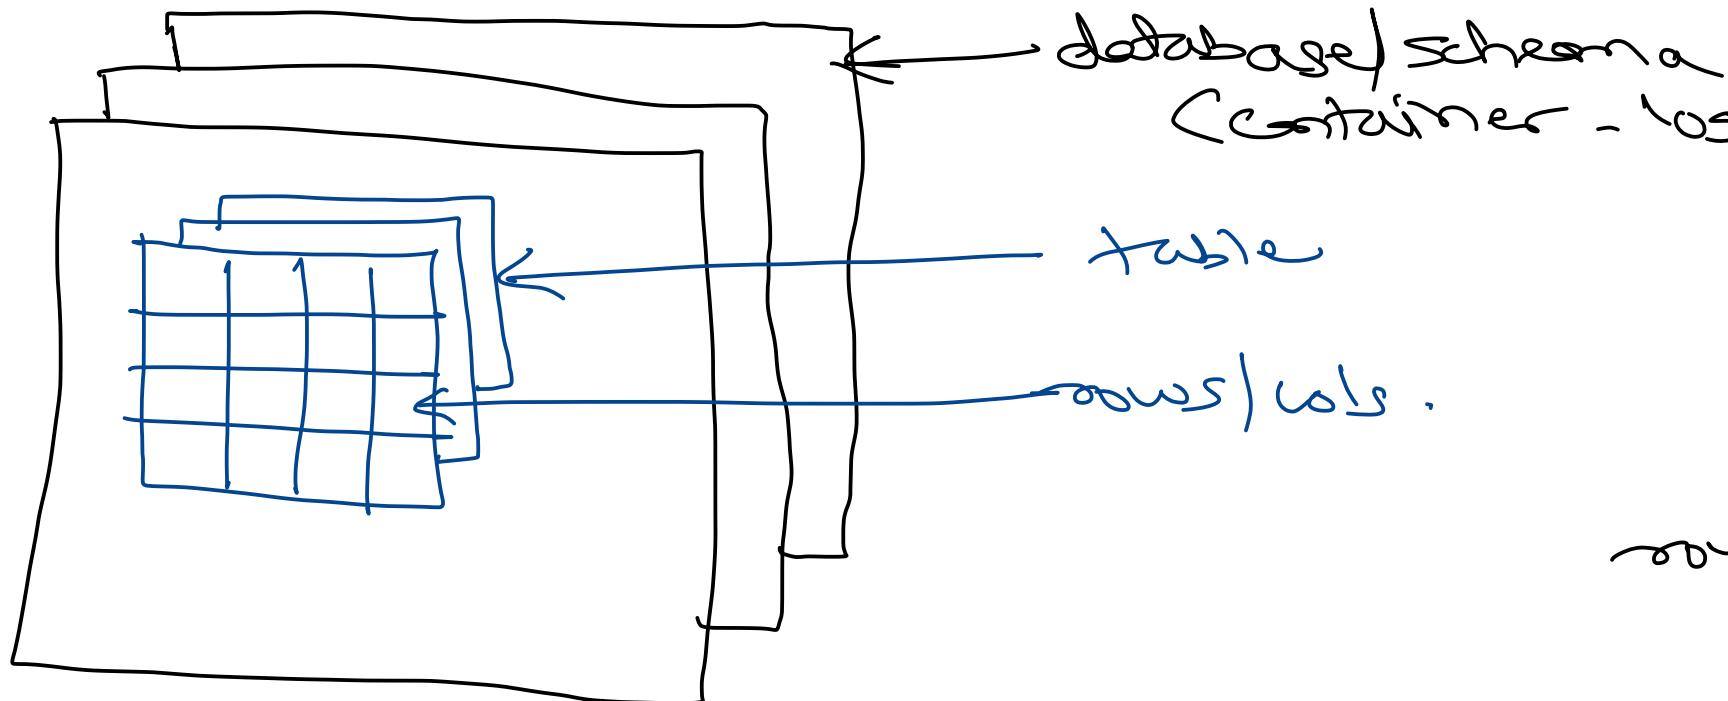
- Database/schema is like a namespace/container that stores all db objects related to a project.
- It contains tables, constraints, relations, stored procedures, functions, triggers, ...
- There are some system databases e.g. mysql, performance\_schema, information\_schema, sys, ... They contain db internal/system information.
  - e.g. SELECT user, host FROM mysql.user;
- A database contains one or more tables.
- Tables have multiple columns.
- Each column is associated with a data-type.
- Columns may have zero or more constraints.
- The data in table is in multiple rows.
- Each row has multiple values (as per columns).

fields

→ separate lecture



*logical*

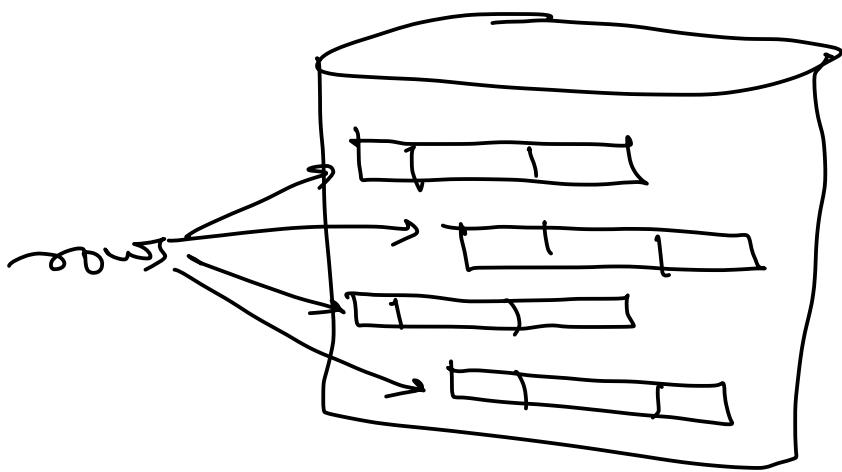


database schema  
(container - logical)

table

rows | cols .

*physical*



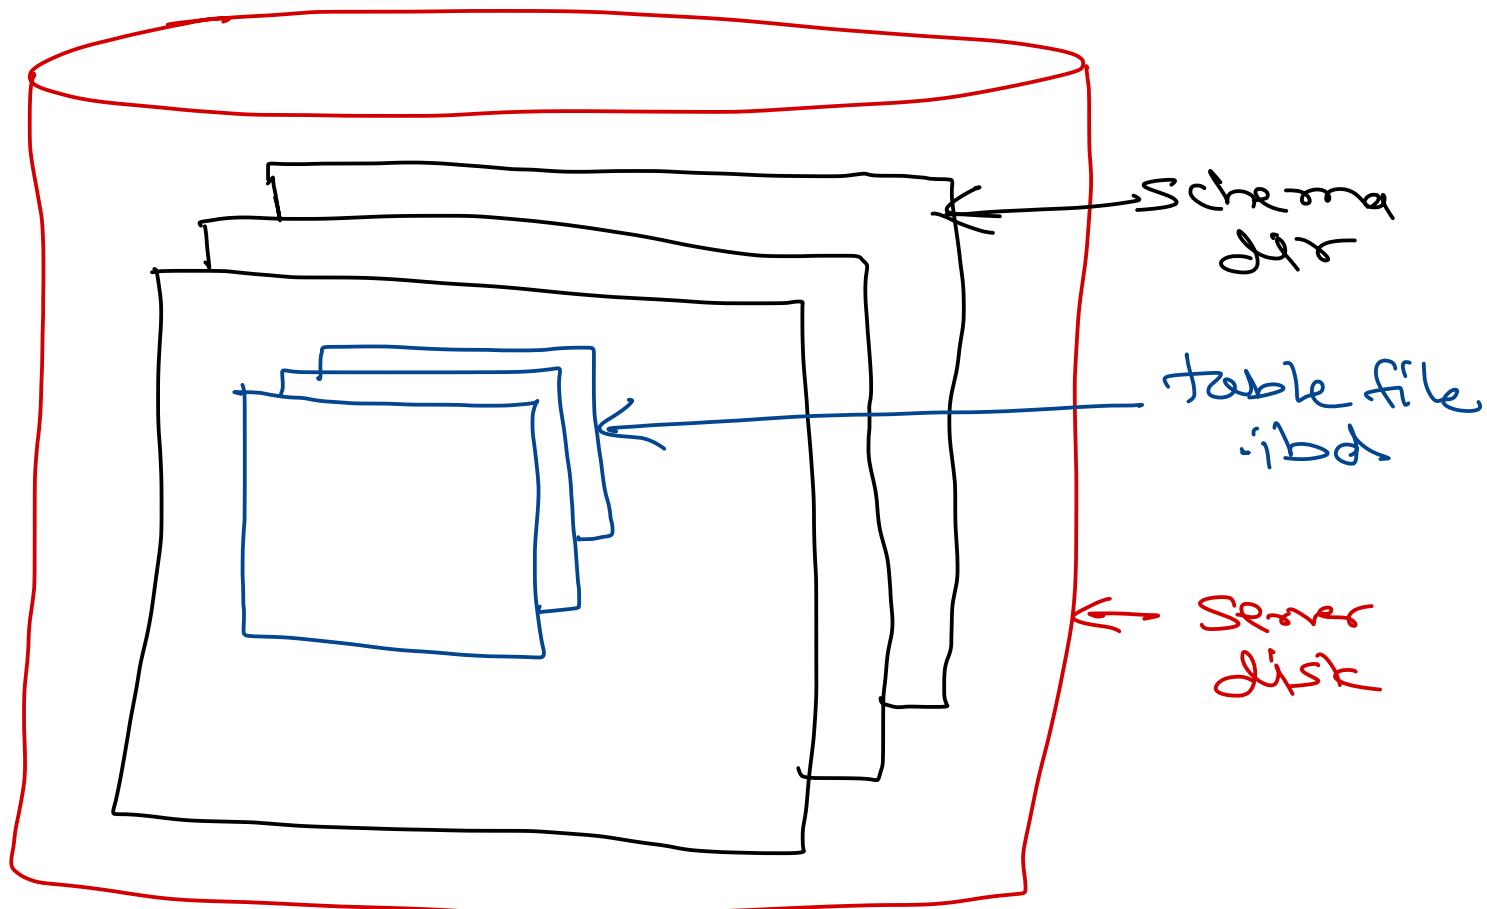
# Database physical layout

Ubuntu

- In MySQL, the data is stored on disk in its data directory i.e. /var/lib/mysql
- Each database/schema is a separate sub-directory in data dir.
- Each table in the db, is a file on disk. (.ibd)
- e.g. student table in current db is stored in file /var/lib/mysql/db/student.ibd.
- Data is stored in binary format.
- A file may not be contiguously stored on hard disk.  
*(grows, file growth: disk alloc)*
- Data rows are not contiguous. They are scattered in the hard disk.
- In one row, all fields are consecutive.
- When records are selected, they are selected in any order.

Select \* from tablename;







# Thank you!

Nilesh Ghule <[nilesh@sunbeaminfo.com](mailto:nilesh@sunbeaminfo.com)>





# RDBMS & SQL Introduction

Trainer: Mr. Nilesh Ghule



# MySQL data types

- RDBMS have similar data types (but not same).

- MySQL data types can be categorised as follows

- Numeric types (Integers)

- TINYINT (1 byte), SMALLINT (2 byte), MEDIUMINT (3 byte), INT (4 byte), BIGINT (8 byte), BIT(n bits)
- integer types can signed (default) or unsigned.

- Numeric types (Floating point)

- approx. precision – FLOAT (4 byte), DOUBLE (8 byte)

- Date/Time types

- DATE, TIME, DATETIME, TIMESTAMP, YEAR

- String types – size = number of chars \* size of char

- CHAR(1-255) – Fixed length, Very fast access.

- VARCHAR(1-65535) – Variable length, Stores length + chars.

- TINYTEXT (255), TEXT (64K), MEDIUMTEXT (16M), LONGTEXT (4G) – Variable length, Slower access.

- Binary types – size = number of bytes

- BINARY, VARBINARY, TINYBLOB, BLOB, MEDIUMBLOB, LONGBLOB

- Miscellaneous types

- ENUM, SET

TINYINT (1)  $\leftrightarrow \pm 2^7$  (signed)  
8 bits.  $\underline{2^8}$  (unsigned).

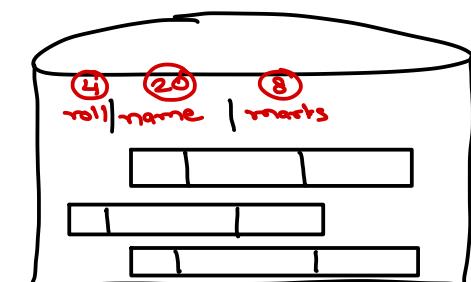
SMALLINT (2)  $\leftrightarrow \pm 2^{15}$   
16 bits  $\underline{2^{16}}$

$\pm 2^{31}$  or  $2^{32}$   $\pm 2^{63}$  or  $2^{64}$

total places/digits.

upto n decimal places

DECIMAL(m, n) – exact precision



# CHAR vs VARCHAR vs TEXT

String {data must be in '\_\_\_\_\_';  
date/time}

- CHAR

- ✓ Fixed inline storage.
- ✓ If smaller data is given, rest of space is unused.
- ✓ Very fast access.

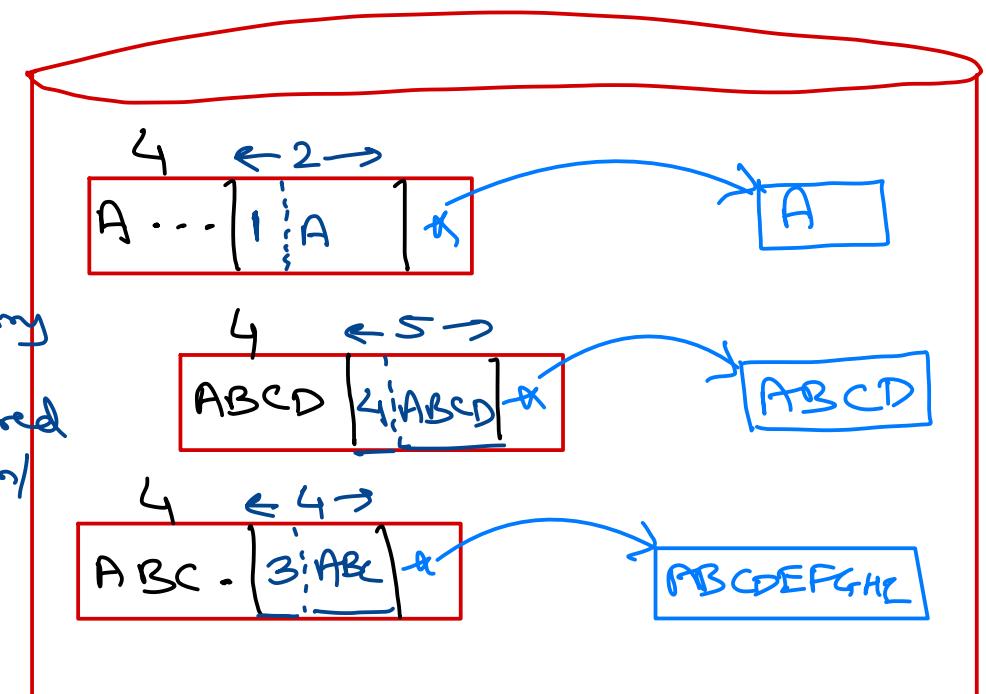
- VARCHAR

- ✓ Variable inline storage.
- ✓ Stores length and characters.
- ✓ Slower access than CHAR.

- TEXT

- ✓ Variable external storage. *to secnd.*
- ✓ Very slow access.
- ✓ Not ideal for indexing. *we will learn in indexing lecture.*

char encoding  
each char → how many bytes  
how stored in mem/disk?



Server disk

Numerical data is not enclosed in '\_\_\_\_\_!'.

✓ CREATE TABLE temp(c1 CHAR(4), c2 VARCHAR(4), c3 TEXT(4));

✓ DESC temp;

✓ INSERT INTO temp VALUES('abcd', 'abcd', 'abcdef');

# SQL scripts

- SQL script is multiple SQL queries written into a .sql file.
  - SQL scripts are mainly used while database backup and restore operations.
  - SQL scripts can be executed from terminal as:
    - terminal> mysql –u user –password db < /path/to/sqlfile
  - SQL scripts can be executed from command line as:
    - mysql> SOURCE /path/to/sqlfile
  - Note that SOURCE is MySQL CLI client command.
  - It reads commands one by one from the script and execute them on server.

each group





# Thank you!

Nilesh Ghule <[nilesh@sunbeaminfo.com](mailto:nilesh@sunbeaminfo.com)>





# RDBMS & SQL Introduction

Trainer: Mr. Nilesh Ghule

# INSERT – DML

- Insert a new row (all columns, fixed order).
  - `INSERT INTO table VALUES (v1, v2, v3);`
- Insert a new row (specific columns, arbitrary order).
  - `INSERT INTO table(c3, c1, c2) VALUES (v3, v1, v2);`
  - `INSERT INTO table(c1, c2) VALUES (v1, v2);`
  - Missing columns data is NULL.
  - NULL is special value and it is not stored in database.
- Insert multiple rows.
  - `INSERT INTO table VALUES (av1, av2, av3), (bv1, bv2, bv3), (cv1, cv2, cv3).`
- Insert rows from another table.
  - `INSERT INTO table SELECT c1, c2, c3 FROM another-table;`
  - `INSERT INTO table (c1,c2) SELECT c1, c2 FROM another-table;`

James	m	1970-1-1
John	m	x

# SELECT – DQL

→ Column order as of  
Creating table.

- Select all columns (in fixed order).

- SELECT \* FROM table;

- Select specific columns / in arbitrary order. → projection .

- SELECT c1, c2, c3 FROM table;

- Column alias

- SELECT c1 AS col1 c2 col2 FROM table;

- Computed columns.

- SELECT c1, c2, c3, expr1, expr2 FROM table;

- SELECT c1,

- CASE WHEN condition1 THEN value1,

- CASE WHEN condition2 THEN value2,

- ...

- ELSE value

- END

- FROM table;

Sal → Salary → 100.0  
Sal → Salutation → Mr. Mrs.

→ Sal or Salary "AS" is optional.

Comm Commission





# Thank you!

Nilesh Ghule <[nilesh@sunbeaminfo.com](mailto:nilesh@sunbeaminfo.com)>



# RDBMS & SQL Introduction

Trainer: Mr. Nilesh Ghule



# SELECT – DQL

- Distinct values in column. → unique values.
  - SELECT DISTINCT c1 FROM table;
  - SELECT DISTINCT c1, c2 FROM table;
- Select limited rows.
  - SELECT \* FROM table LIMIT n;
  - SELECT \* FROM table LIMIT m, n;

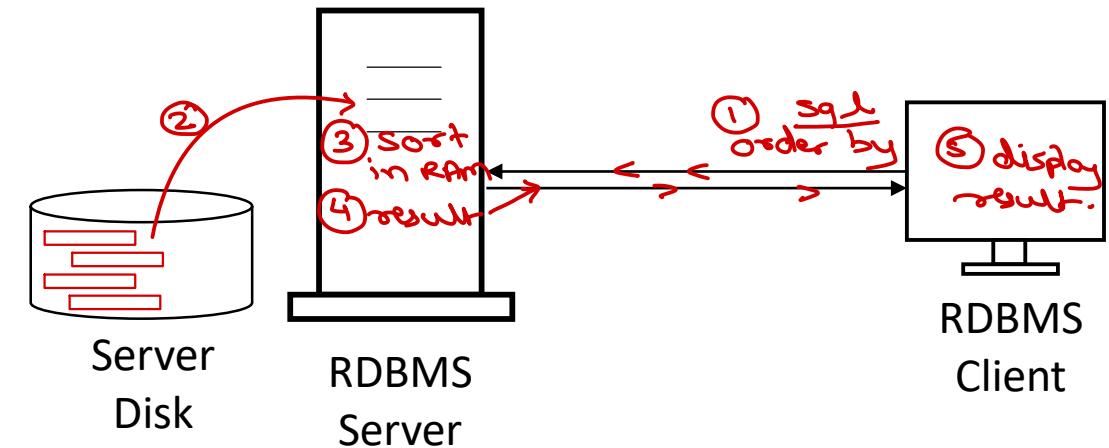


# SELECT – DQL – ORDER BY

→ Sorting  asc (default)  
                  ↓  
                  desc

order by sal desc;  
desc emp;  
↳ describe table struct

- In db rows are scattered on disk. Hence may not be fetched in a fixed order.
  - Select rows in asc order.
    - `SELECT * FROM table ORDER BY c1;`
    - `SELECT * FROM table ORDER BY c2 ASC;`
  - Select rows in desc order.
    - `SELECT * FROM table ORDER BY c3 DESC;`
  - Select rows sorted on multiple columns.
    - `SELECT * FROM table ORDER BY c1, c2;`
    - `SELECT * FROM table ORDER BY c1 ASC, c2 DESC;`
    - `SELECT * FROM table ORDER BY c1 DESC, c2 DESC;`
  - Select top or bottom n rows.
    - `SELECT * FROM table ORDER BY c1 ASC LIMIT n;`
    - `SELECT * FROM table ORDER BY c1 DESC LIMIT n;`
    - `SELECT * FROM table ORDER BY c1 ASC LIMIT m, n;`



# SELECT – DQL – WHERE

- It is always good idea to fetch only required rows (to reduce network traffic).
- The WHERE clause is used to specify the condition, which records to be fetched.
- ✓ Relational operators
  - <, >, <=, >=, =, != or  $\diamond$
- ✓ NULL related operators
  - NULL is special value and cannot be compared using relational operators.
  - IS NULL or <=>, IS NOT NULL.
- Logical operators
  - AND, OR, NOT

operators

① between  
② in  
③ like

} monday  
lectures .





# Thank you!

Nilesh Ghule <[nilesh@sunbeaminfo.com](mailto:nilesh@sunbeaminfo.com)>





# RDBMS & SQL Introduction

Trainer: Mr. Nilesh Ghule



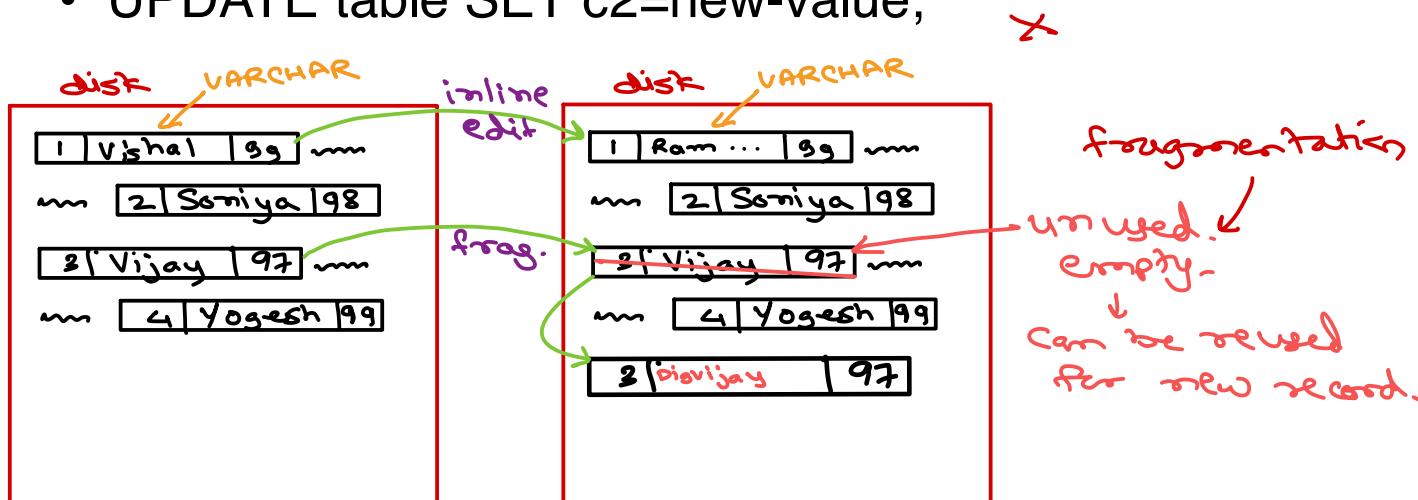
# SELECT – DQL – WHERE

- BETWEEN operator (include both ends) NOT BETWEEN
  - ✓ • c1 BETWEEN val1 AND val2
- IN operator (equality check with multiple values) NOT IN
  - ✓ • c1 IN (val1, val2, val3)
- LIKE operator (similar strings) NOT LIKE
  - ✓ • c1 LIKE 'pattern'.
  - ✓ • % represent any number of any characters.
  - ✓ • \_ represent any single character.



# UPDATE – DML

- To change one or more rows in a table.
- Update row(s) single column.
  - UPDATE table SET c2=new-value WHERE c1=some-value;
- Update multiple columns.
  - UPDATE table SET c2=new-value, c3=new-value WHERE c1=some-value;
- Update all rows single column.
  - UPDATE table SET c2=new-value;



# DELETE – DML vs TRUNCATE – DDL vs DROP – DDL

## • DELETE

- To delete one or more rows in a table.
- Delete row(s)
  - DELETE FROM table WHERE c1=value;
- Delete all rows
  - DELETE FROM table ;

*not columns*

*↓*

*✓*



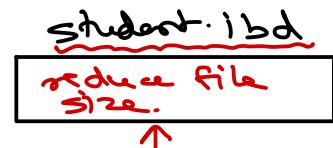
← Students  
table

DELETE ...  
→ mark as deleted.  
- fragmentation.

← Can be  
rolled back  
in  
transaction.

## • TRUNCATE

- Delete all rows.
  - TRUNCATE TABLE table;
- Truncate is faster than DELETE.

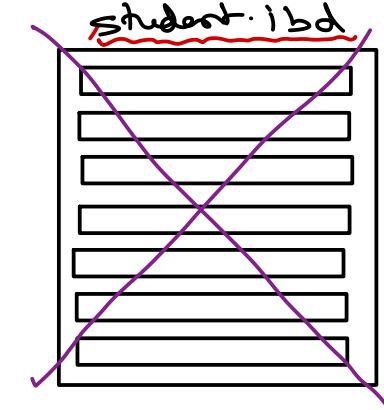


TRUNCATE

↓  
faster than delete all.  
cannot be rolled back.

## • DROP

- Delete all rows as well as table structure.
  - DROP TABLE table;
  - DROP TABLE table IF EXISTS;
- Delete database/schema.
  - DROP DATABASE db; → root login



← Drop.  
↳ faster





# Thank you!

Nilesh Ghule <[nilesh@sunbeaminfo.com](mailto:nilesh@sunbeaminfo.com)>



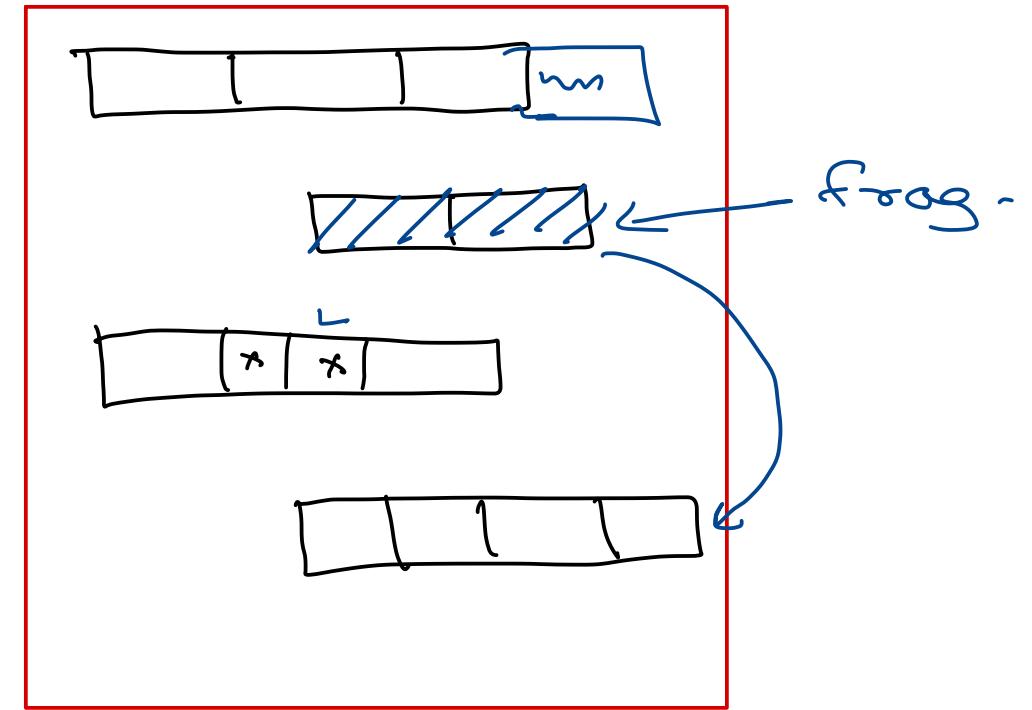
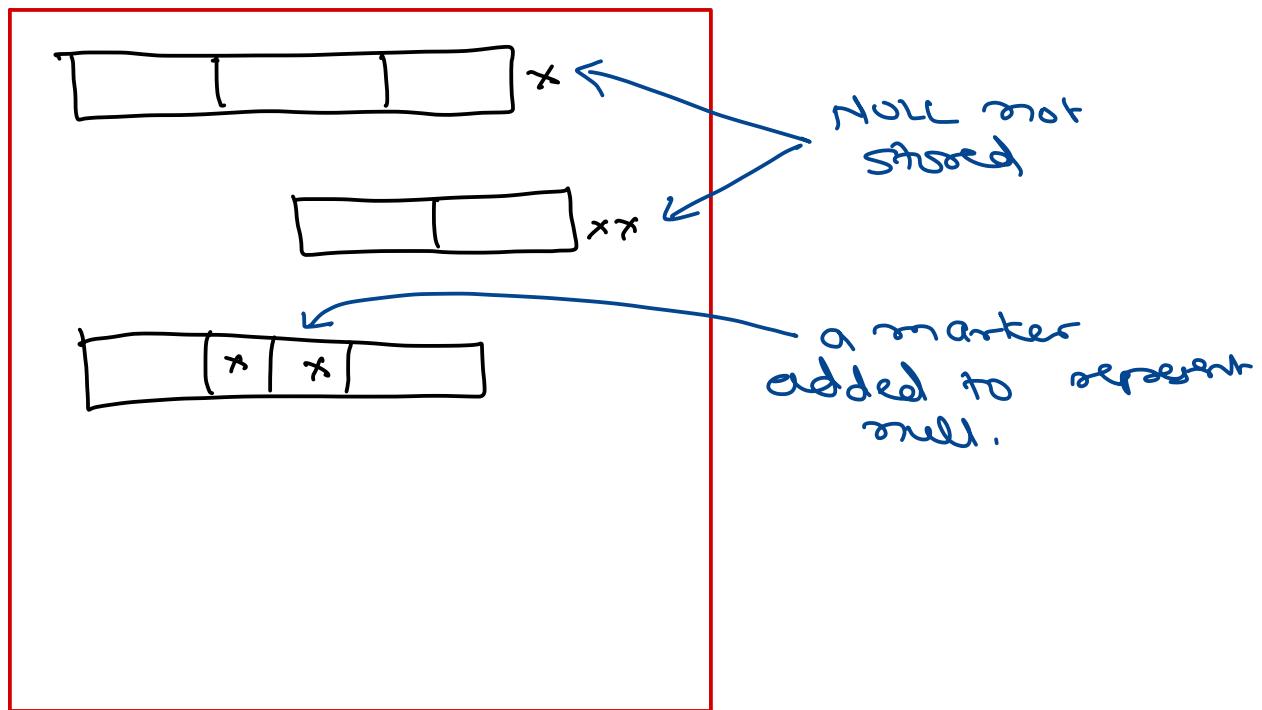


# MySQL RDBMS

Trainer: Mr. Nilesh Ghule



NULL



# Transaction

- Transaction is set of DML queries executed as a single unit.
- Transaction examples
  - accounts table [id, type, balance]
  - + ✓ • UPDATE accounts SET balance=balance-1000 WHERE id = 1;
  - ✗ ✗ • UPDATE accounts SET balance=balance+1000 WHERE id = 2;
- RDBMS transaction have ACID properties.
  - Atomicity
    - All queries are executed as a single unit. If any query is failed, other queries are discarded.
  - Consistency
    - When transaction is completed, all clients see the same data.
  - Isolation
    - Multiple transactions (by same or multiple clients) are processed concurrently.
  - Durable
    - When transaction is completed, all data is saved on disk.



# Transaction

- Transaction management
  - START TRANSACTION;
  - ... ~~≡~~
  - COMMIT WORK;
- START TRANSACTION;  
• ... ~~≡~~  
• ROLLBACK WORK;
- In MySQL autocommit variable is by default 1. So each DML command is auto-committed into database. → each query is one tx - auto committed.
  - SELECT @@autocommit;
- Changing autocommit to 0, will create new transaction immediately after current transaction is completed. This setting can be made permanent in config file.
  - SET autocommit=0;



# Transaction

- Save-point is state of database tables (data) at the moment (within a transaction).
- It is advised to create save-points at end of each logical section of work.
- Database user may choose to rollback to any of the save-point.
- Transaction management with Save-points
  - START TRANSACTION;
  - ... =
  - SAVEPOINT sa1; ✓ ↗
  - ... =
  - SAVEPOINT sa2; ↗
  - ... =
  - ROLLBACK TO sa1; ↗
  - ... =
  - COMMIT; // or ROLLBACK
- Commit always commit the whole transaction.
- ROLLBACK or COMMIT clears all save-points.

```
start transaction;  
✓ { dml1;  
    dml2;  
    savepoint sa1; ↗  
    dml3; X  
    dml4;  
    savepoint sa2; ↗  
    dml5; X  
    dml6; X  
    rollback to sa1;  
✓ { dml5;  
    dml6;  
    commit;
```



# Transaction

---

- Transaction is set of DML statements.
  - If any DDL statement is executed, current transaction is automatically committed.
  - Any power failure, system or network failure automatically rollback current state.
  - Transactions are isolated from each other and are consistent.
- 



# Row locking

- When an user update or delete a row (within a transaction), that row is locked and becomes read-only for other users.
- The other users see old row values, until transaction is committed by first user.
- If other users try to modify or delete such locked row, their transaction processing is blocked until row is unlocked.
- Other users can INSERT into that table. Also they can UPDATE or DELETE other rows.
- The locks are automatically released when COMMIT/ROLLBACK is done by the user.
- This whole process is done automatically in MySQL. It is called as "OPTIMISTIC LOCKING".



# Row locking

- Manually locking the row in advance before issuing UPDATE or DELETE is known as "PESSIMISTIC LOCKING".
- This is done by appending FOR UPDATE to the SELECT query.
- It will lock all selected rows, until transaction is committed or rolled back.
- If these rows are already locked by another users, the SELECT operation is blocked until rows lock is released.
- By default MySQL does table locking. Row locking is possible only when table is indexed on the column.





# Thank you!

Nilesh Ghule <[nilesh@sunbeaminfo.com](mailto:nilesh@sunbeaminfo.com)>





# MySQL RDBMS

Trainer: Mr. Nilesh Ghule



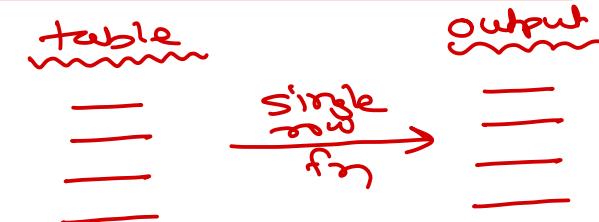
# SQL functions

- RDBMS provides many built-in functions to process the data.

- These functions can be classified as:

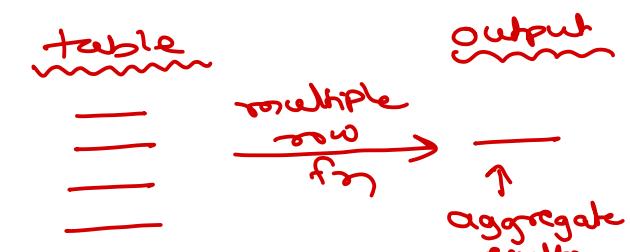
- Single row functions

- One row input produce one row output.
    - e.g. ABS(), CONCAT(), IFNULL(), ...



- Multi-row or Group functions

- Values from multiple rows are aggregated to single value.
    - e.g. SUM(), MIN(), MAX(), ...



- These functions can also be categorized based on data types or usage.

- Numeric functions
  - String functions
  - Date and Time functions
  - Control flow functions
  - Information functions
  - Miscellaneous functions

# Numeric & String functions

- ABS()
- POWER()
- ROUND(), FLOOR(), CEIL()
  
- ASCII(), CHAR()
- CONCAT()
- SUBSTRING()
- LOWER(), UPPER()
- TRIM(), LTRIM(), RTRIM()
- LPAD(), RPAD()
- REGEXP\_LIKE()



# Date-Time and Information functions

- VERSION()
- USER(), DATABASE()
- MySQL supports multiple date time related data types
  - DATE (3), TIME (3), DATETIME (5), TIMESTAMP (4), YEAR (1)
- SYSDATE(), NOW()
- DATE(), TIME()
- DAYOFMONTH(), MONTH(), YEAR(), HOUR(), MINUTE(), SECOND(), ...
- DATEDIFF(), DATE\_ADD(), TIMEDIFF()
- MAKEDATE(), MAKETIME()



# Control and NULL and List functions

---

- NULL is special value in RDBMS that represents absence of value in that column.
  - NULL values do not work with relational operators and need to use special operators.
  - Most of functions return NULL if NULL value is passed as one of its argument.
  - ISNULL()
  - IFNULL()
  - NULLIF()
  - COALESCE()
- 
- GREATEST(), LEAST()
  - IF(condition, true-value, false-value)





# Thank you!

Nilesh Ghule <[nilesh@sunbeaminfo.com](mailto:nilesh@sunbeaminfo.com)>





# MySQL RDBMS

Trainer: Mr. Nilesh Ghule



# Control and NULL and List functions

---

- NULL is special value in RDBMS that represents absence of value in that column.
  - NULL values do not work with relational operators and need to use special operators.
  - Most of functions return NULL if NULL value is passed as one of its argument.
  - ISNULL()
  - IFNULL()
  - NULLIF()
  - COALESCE()
- 
- GREATEST(), LEAST()
- 
- IF(condition, true-value, false-value)



# Group functions

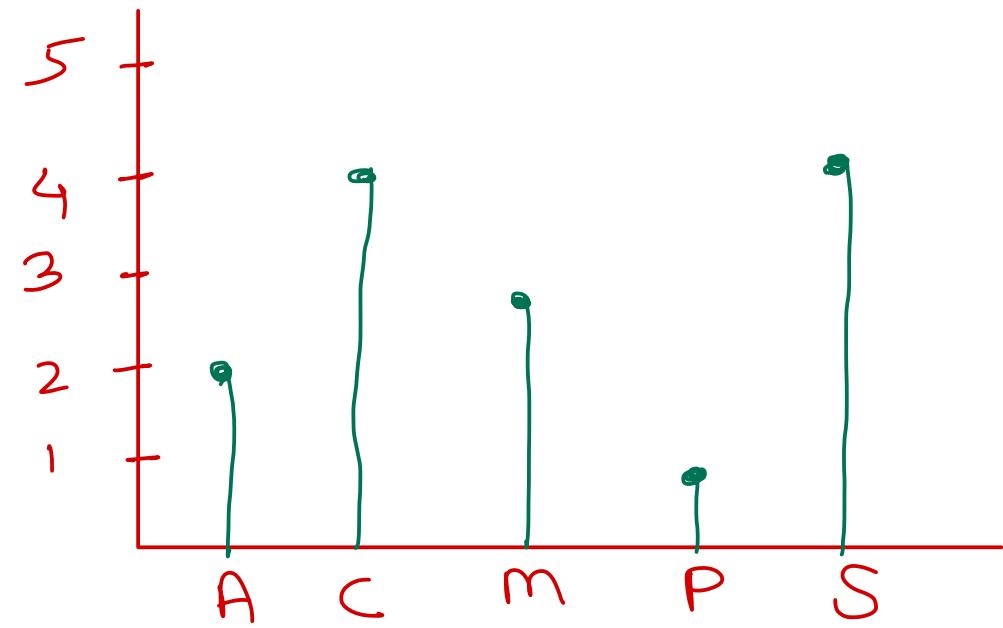
- Work on group of rows of table.
- Input to function is data from multiple rows & then output is single row. Hence these functions are called as "Multi Row Function" or "Group Functions".
- These functions are used to perform aggregate ops like sum, avg, max, min, count or std dev, etc. Hence these fns are also called as "Aggregate Functions".
- Example: SUM(), AVG(), MAX(), MIN(), COUNT().
- NULL values are ignored by group functions.
- Limitations of GROUP functions:
  - Cannot select group function along with a column.
  - Cannot select group function along with a single row fn.
  - Cannot use group function in WHERE clause/condition.
  - Cannot nest a group function in another group fn.



## GROUP BY clause

- GROUP BY is used for analysis of data i.e. generating reports & charts.
- When GROUP BY single column, generated output can be used to plot 2-D chart.  
When GROUP BY two column, generated output can be used to plot 3-D chart and so on.
- GROUP BY queries are also called as Multi-dimensional / Spatial queries.
- Syntactical Characteristics:
  - If a column is used for GROUP BY, then it may or may not be used in SELECT clause.
  - If a column is in SELECT, it must be in GROUP BY.
- When GROUP BY query is fired on database server, it does following:
  - Load data from server disk into server RAM.
  - Sort data on group by columns.
  - Group similar records by group columns.
  - Perform given aggregate ops on each column.
  - Send result to client.







# Thank you!

Nilesh Ghule <[nilesh@sunbeaminfo.com](mailto:nilesh@sunbeaminfo.com)>





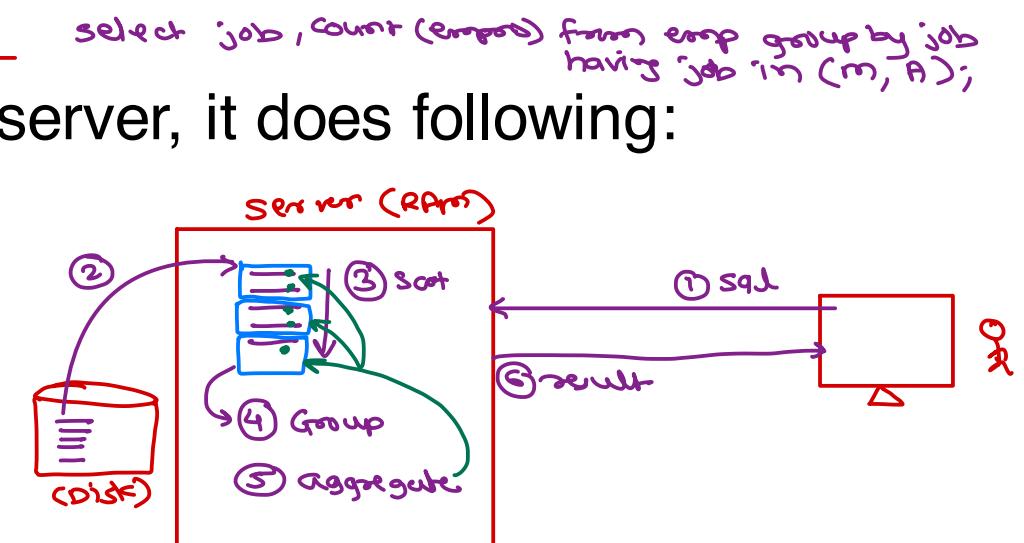
# MySQL RDBMS

Trainer: Mr. Nilesh Ghule



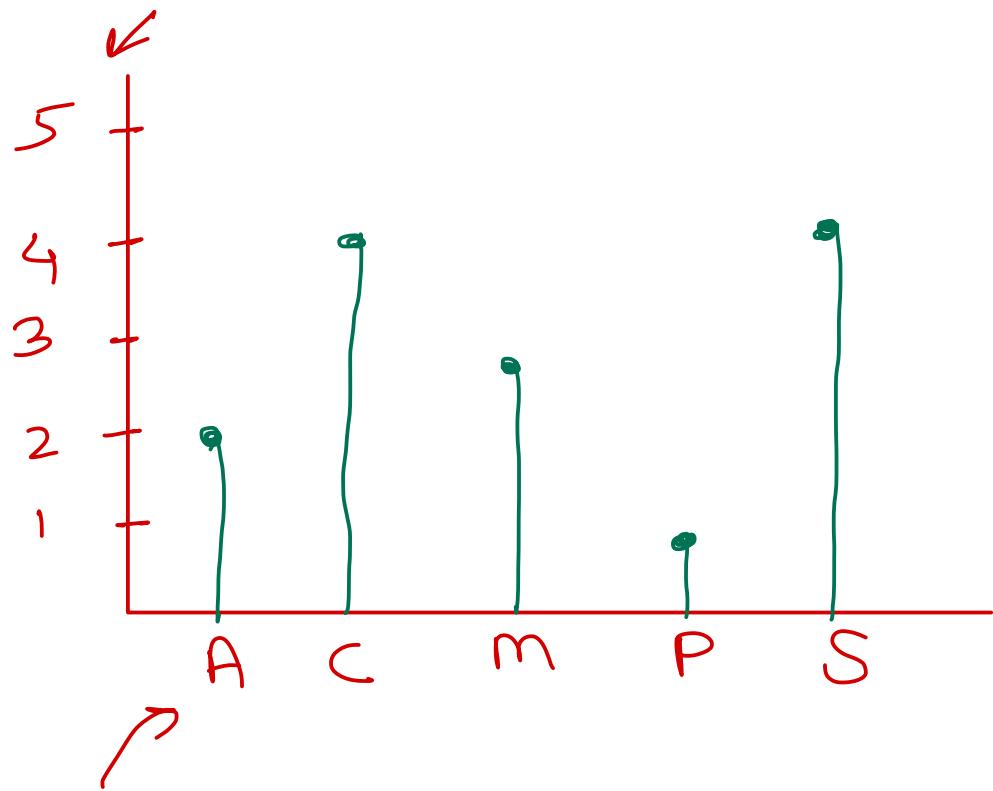
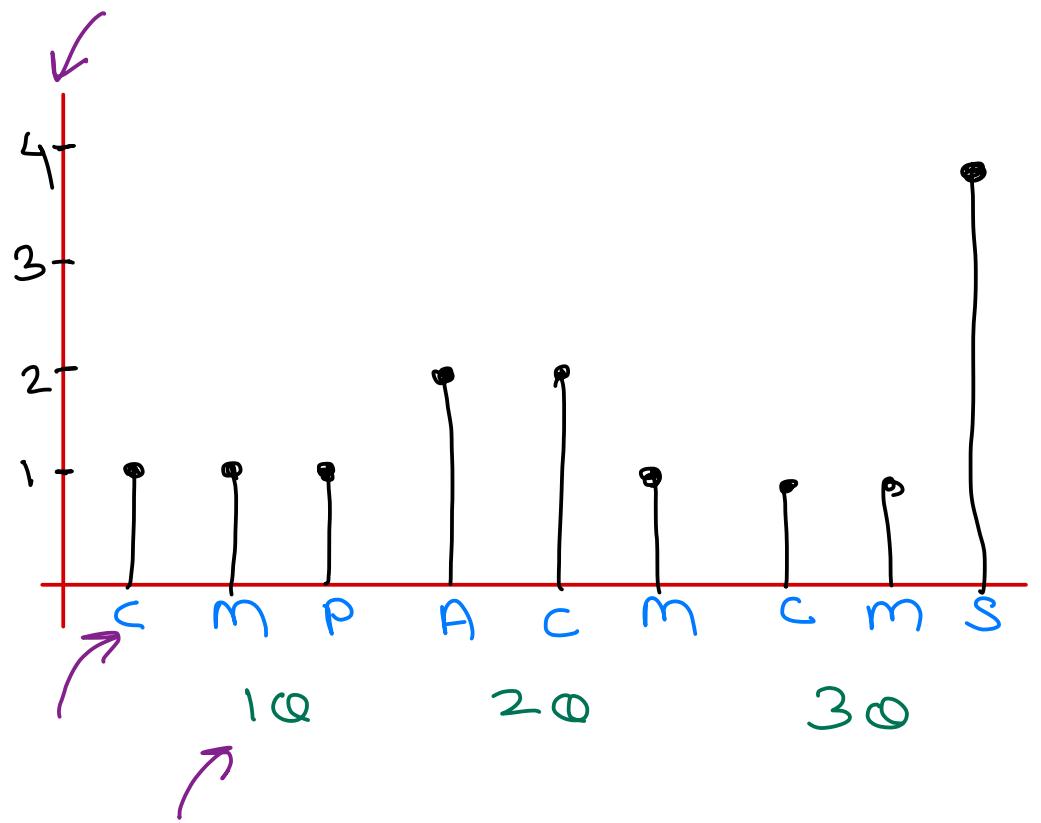
# GROUP BY clause

- GROUP BY is used for analysis of data i.e. generating reports & charts.
- When GROUP BY single column, generated output can be used to plot 2-D chart.  
When GROUP BY two column, generated output can be used to plot 3-D chart and so on.
- GROUP BY queries are also called as Multi-dimensional / Spatial queries.
- Syntactical Characteristics:
  - If a column is used for GROUP BY, then it may or may not be used in SELECT clause.
  - If a column is in SELECT, it must be in GROUP BY.
- When GROUP BY query is fired on database server, it does following:
  - ✓ Load data from server disk into server RAM.
  - ✓ Sort data on group by columns.
  - ✓ Group similar records by group columns.
  - ✓ Perform given aggregate ops on each column.
  - ✓ Send result to client.



select job, count(emps) from emp where job in ('M', 'A')  
group by job  
www.sunbeaminfo.com





# HAVING clause

- HAVING clause cannot be used without GROUP BY clause.
- HAVING clause is used to specify condition on aggregate values *fns*.
- Examples:
  - SELECT deptno, SUM(sal) FROM EMP GROUP BY deptno HAVING SUM(sal) > 9000;
- Syntactical Characteristics:
  - WHERE clause executed for each record; while HAVING is executed for each group.
  - HAVING clause can be used to specify condition on group fn or grouped columns.
  - However using HAVING to specify condition of group col reduce the performance. Use WHERE clause for the same.
- Examples:
  - ✓ SELECT deptno, SUM(sal) FROM EMP GROUP BY deptno HAVING deptno = 20;
  - ✗ SELECT deptno, SUM(sal) FROM EMP WHERE deptno = 20 GROUP BY deptno; *efficient.*
- We may use GROUP BY with WHERE, ORDER BY & LIMIT.





# Thank you!

Nilesh Ghule <[nilesh@sunbeaminfo.com](mailto:nilesh@sunbeaminfo.com)>





# MySQL RDBMS

Trainer: Mr. Nilesh Ghule



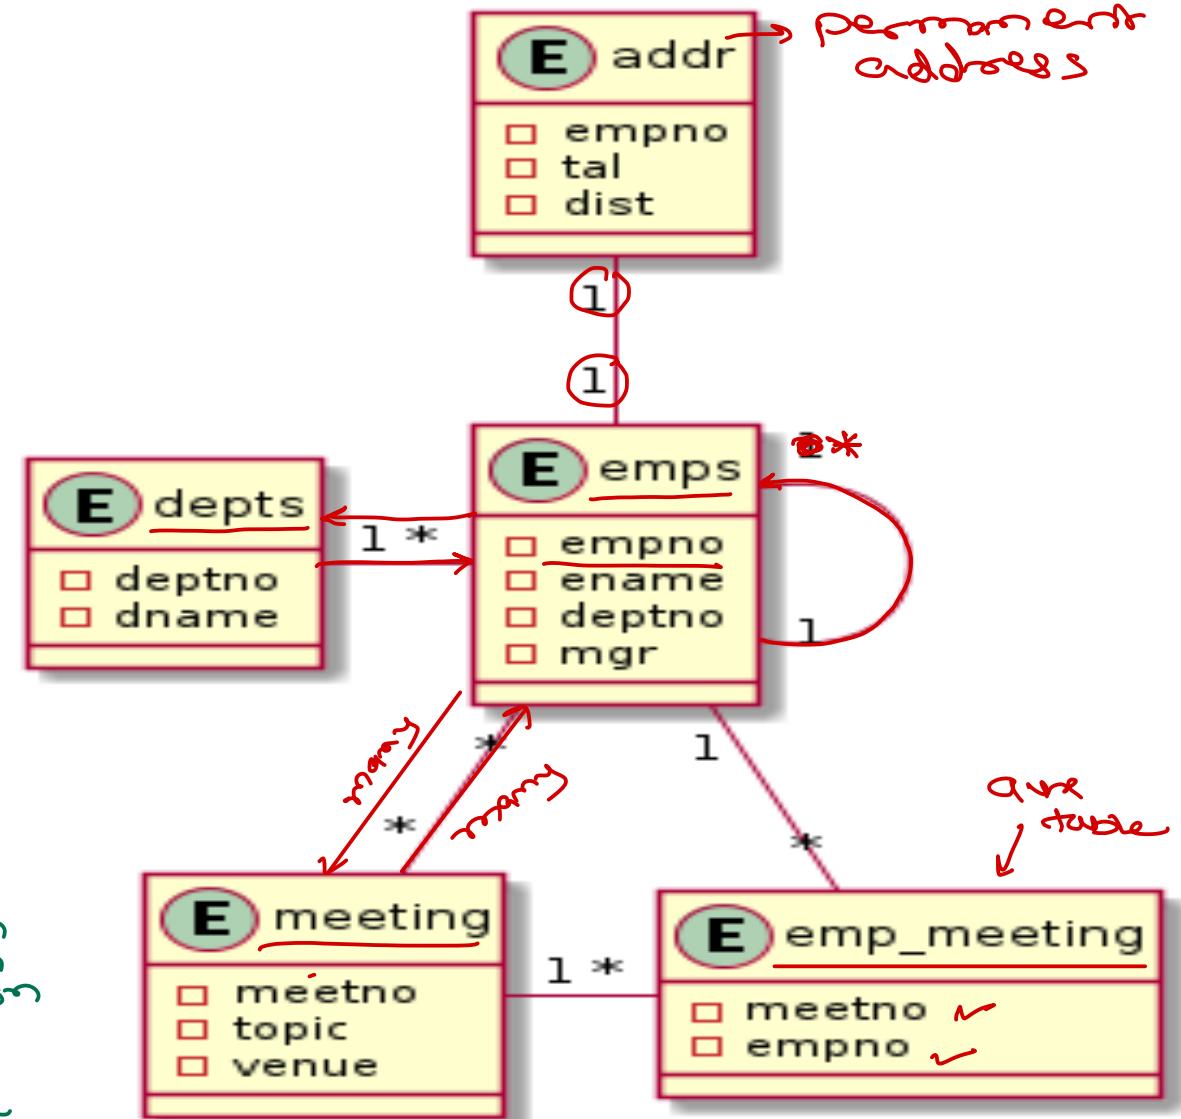
# Entity Relations

## ER - diagram

- To avoid redundancy of the data, data should be organized into multiple tables so that tables are related to each other.
- The relations can be one of the following
  - One to One ✓
  - One to Many ✓
  - Many to One ✓
  - Many to Many ✓
- Entity relations is outcome of Normalization process.

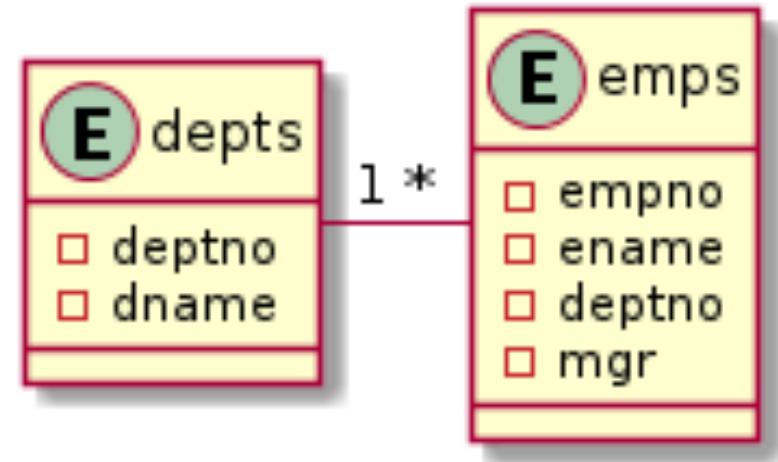
part of application design

- ① req analysis → ui design
- ② design ← db design → oo design
- ③ implementation
- ④ testing / maintenance



# SQL Joins

- Join statements are used to **SELECT** data from **multiple tables** using **single query**.
- Typical RDBMS supports following types of joins:
  - ✓ Cross Join
  - ✓ Inner Join
  - ✓ Left Outer Join
  - ✓ Right Outer Join
  - ✓ Full Outer Join
  - ✓ Self join



# Cross Join

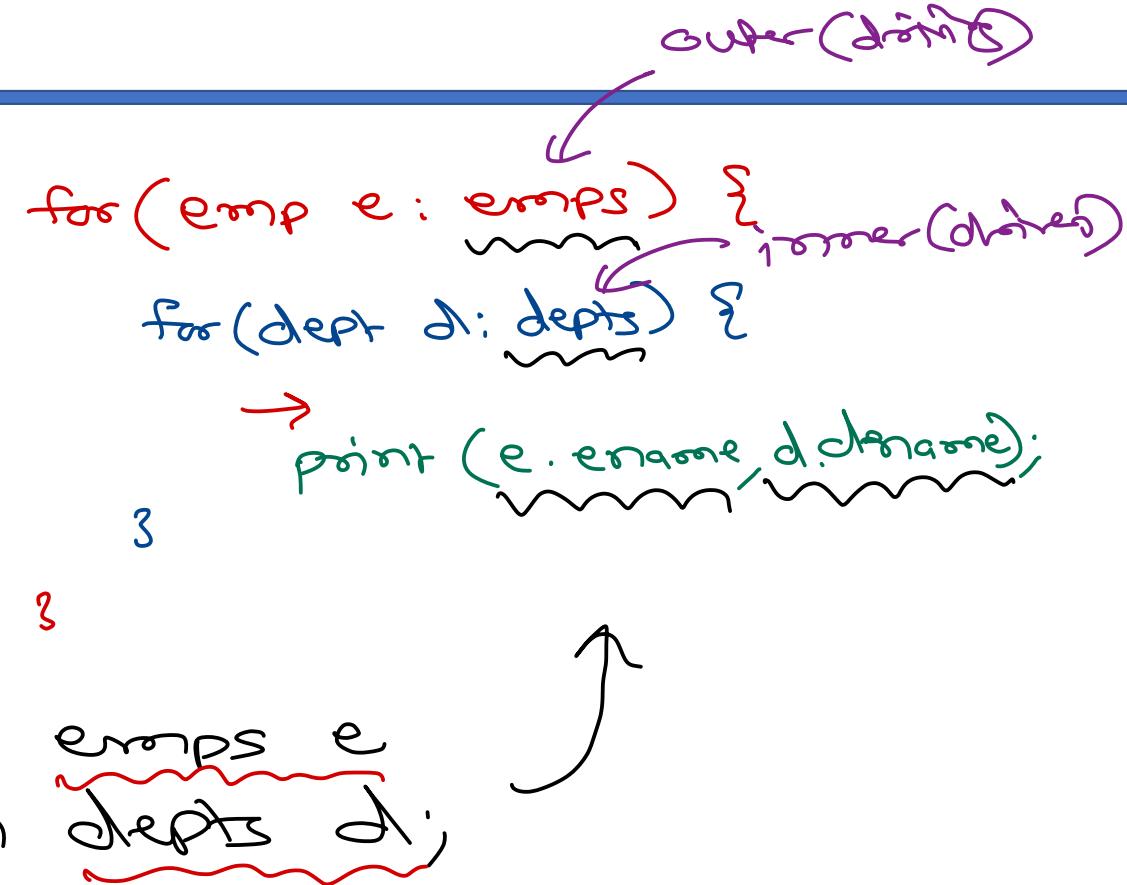
depts (4)      emps (5)

deptno	dname
10	DEV
20	QA
30	OPS
40	ACC

empno	ename	deptno
1	Amit	10
2	Rahul	10
3	Nilesh	20
4	Nitin	50
5	Sarang	50

Select e.ename, d.dname from  
cross join



# Inner Join → common rows

deptno	dname
10	DEV
20	QA
30	OPS
40	ACC

empno	ename	deptno
1	Amit	10
2	Rahul	10
3	Nilesh	20
4	Nitin	50
5	Sarang	50

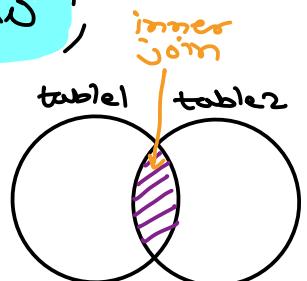
so

select e.ename, d.dname from emps e  
 inner join depts d on e.deptno = d.deptno;

```
for(emp e: emps) {
  for(dept d: depts) {
    if(e.deptno == d.deptno)
      print(e.ename, d.dname);
}
```

3 3

Amit	DEV
Rahul	DEV
Nilesh	QA



- The inner JOIN is used to return rows from both tables that satisfy the join condition.
- Non-matching rows from both tables are skipped.
- If join condition contains equality check, it is referred as equi-join; otherwise it is non-equi-join.



# Thank you!

Nilesh Ghule <[nilesh@sunbeaminfo.com](mailto:nilesh@sunbeaminfo.com)>





# MySQL RDBMS

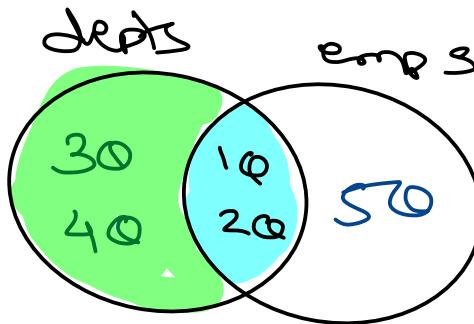
Trainer: Mr. Nilesh Ghule



# Left Outer Join

deptno	dname
10	DEV
20	QA
30	OPS
40	ACC

empno	ename	deptno
1	Amit	10
2	Rahul	10
3	Nilesh	20
4	Nitin	50
5	Sarang	50



```
for (Dept d : depts) {
    found = 0;
    for (Emp e : emps) {
        if (d.deptno == e.deptno) {
            print(e.ename, d.dname);
            found = 1;
        }
    }
    if (found == 0)
        print(NULL, d.dname);
}
```

select e.ename, d.dname from depts d  
left outer join emps e on d.deptno = e.deptno;

- Left outer join is used to return matching rows from both tables along with additional rows in left table.
- Corresponding to additional rows in left table, right table values are taken as NULL.
- OUTER keyword is optional.

# Right Outer Join

deptno	dname
10	DEV
20	QA
30	OPS
40	ACC



empno	ename	deptno
1	Amit	10
2	Rahul	10
3	Nilesh	20
4	Nitin	50
5	Sarang	50

select e.ename, d.dname from depts d  
right outer join emps e on d.deptno = e.deptno;

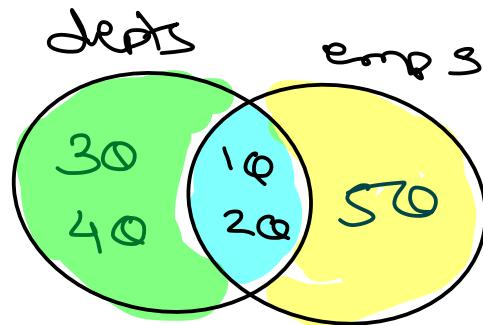
```
for (Emp e : emps) {  
    found = 0;  
    for (Dept d : depts) {  
        if (d.deptno == e.deptno) {  
            print(e.ename, d.dname);  
            found = 1;  
        }  
    }  
    if (found == 0)  
        print(e.ename, null);  
}
```

- Right outer join is used to return matching rows from both tables along with additional rows in right table.
- Corresponding to additional rows in right table, left table values are taken as NULL.
- OUTER keyword is optional.



# Full Outer Join

deptno	dname
10	DEV
20	QA
30	OPS
40	ACC



empno	ename	deptno
1	Amit	10
2	Rahul	10
3	Nilesh	20
4	Nitin	50
5	Sarang	50

## Left Join

Amit	Dev
Rahul	Dev
Nilesh	QA
+	OPS
+	ACC

## Right Join

Amit	Dev
Rahul	Dev
Nilesh	QA
Nitin	X
Sarang	X

## Full Join

Amit	Dev
Rahul	Dev
Nilesh	QA
+	OPS
+	ACC
Nitin	X
Sarang	X

- Full join is used to return matching rows from both tables along with additional rows in both tables.
- Corresponding to additional rows in left or right table, opposite table values are taken as NULL.
- Full outer join is not supported in MySQL, but can be simulated using set operators.



# Set operators – Combine output of two queries.

ename	dname
Amit	DEV
Rahul	DEV
Nilesh	QA
NULL	OPS
NULL	ACC

ename	dname
Amit	DEV
Rahul	DEV
Nilesh	QA
Nitin	NULL
Sarang	NULL

Left Join

Amit Dev  
Rahul Dev  
Nilesh QA  
+ OPS  
X ACC

Right Join

Amit Dev  
Rahul Dev  
Nilesh QA  
Nitin X  
Sarang X

union all

Amit Dev -  
Rahul Dev -  
Nilesh QA -  
+ OPS  
X ACC  
Amit Dev -  
Rahul Dev -  
Nilesh QA -  
Nitin X  
Sarang X

Union

Amit Dev  
Rahul Dev  
Nilesh QA  
+ OPS  
X ACC  
Nitin X  
Sarang X

like  
full outer  
join  
of  
oracle /  
ms-sql.

- UNION operator is used to combine results of two queries. The common data is taken only once. It can be used to simulate full outer join.
- UNION ALL operator is used to combine results of two queries. Common data is repeated.



# Self Join

- When join is done on same table, then it is known as "Self Join". The both columns in condition belong to the same table.
- Self join may be an inner join or outer join.

empno	ename	deptno	mgr
1	Amit	10	4
2	Rahul	10	3
3	Nilesh	20	4
4	Nitin	50	5
5	Sarang	50	NULL

empno	ename	deptno	mgr
1	Amit	10	4
2	Rahul	10	3
3	Nilesh	20	4
4	Nitin	50	5
5	Sarang	50	NULL

*mgr column - represent empno  
of manager of current emp.*

# Self Join

- When join is done on same table, then it is known as "Self Join". The both columns in condition belong to the same table.
- Self join may be an inner join or outer join.

empno	ename	deptno	mgr
1	Amit	10	4
2	Rahul	10	3
3	Nilesh	20	4
4	Nitin	50	5
5	Sarang	50	NULL

empno	ename	deptno	mgr
1	Amit	10	4
2	Rahul	10	3
3	Nilesh	20	4
4	Nitin	50	5
5	Sarang	50	NULL

e      m  
Amit, Nitin  
Rahul, Nilesh  
Nilesh, Nitin  
Nitin, Sarang

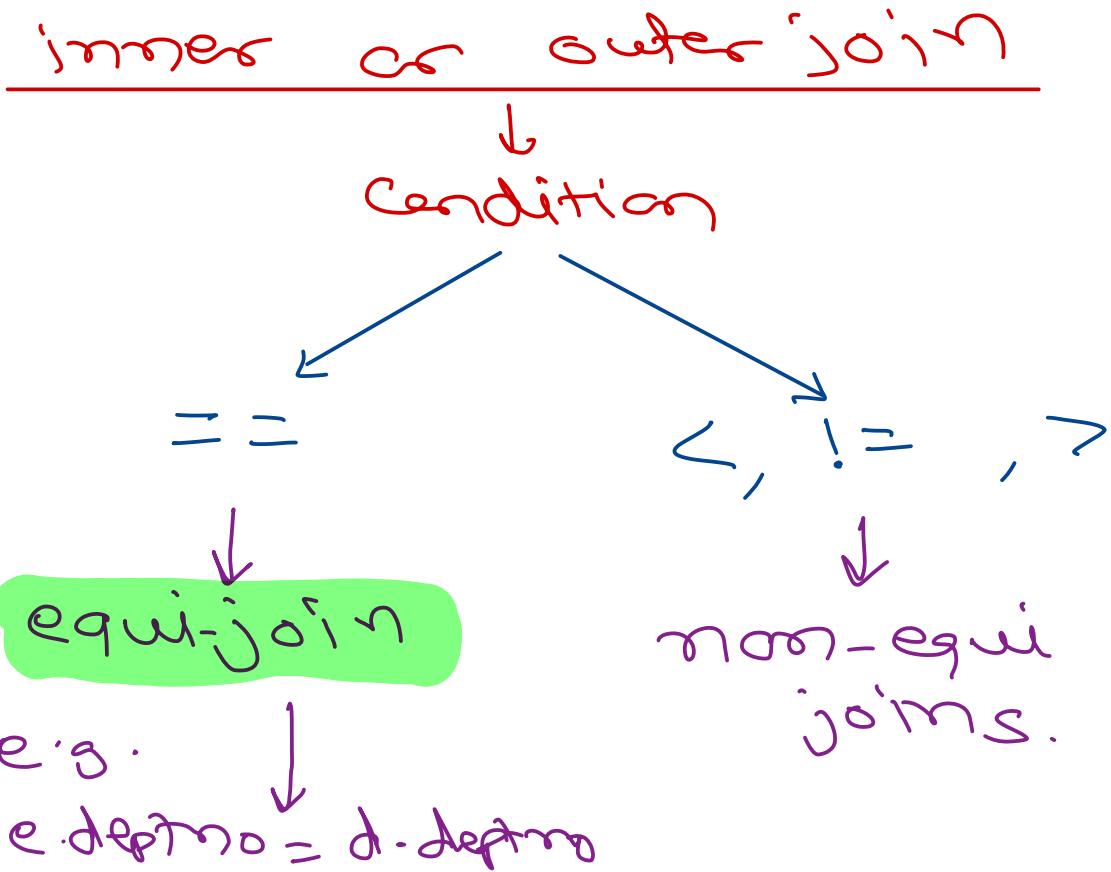
e      m  
Amit, Nitin  
Rahul, Nilesh  
Nilesh, Nitin  
Nitin, Sarang  
Sarang, X

Select e.ename, m.ename from emps e  
inner join emps m on e.mgr = m.empno;

Select e.ename, m.ename from emps e  
left join emps m on e.mgr = m.empno; ↗



# Joins



Select e.empname, d.dname from  
depts d inner join emps e  
on e.deptno = d.deptno; ✓

Select e.empname, d.dname  
from  
depts d join emps e  
using (deptno);

↳ col name is same in  
both tables.

## Joins - Inner join

### Standard

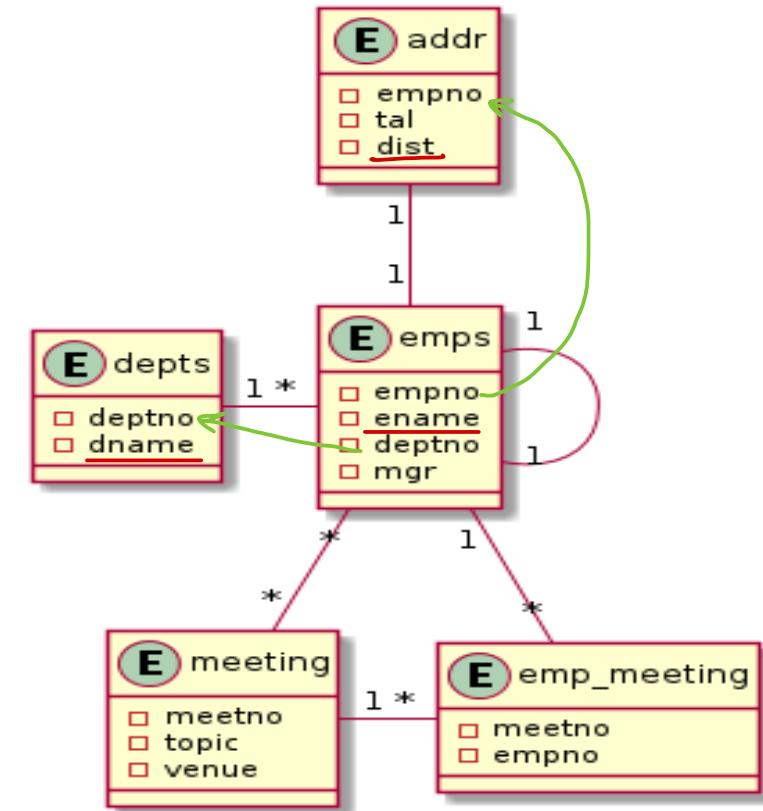
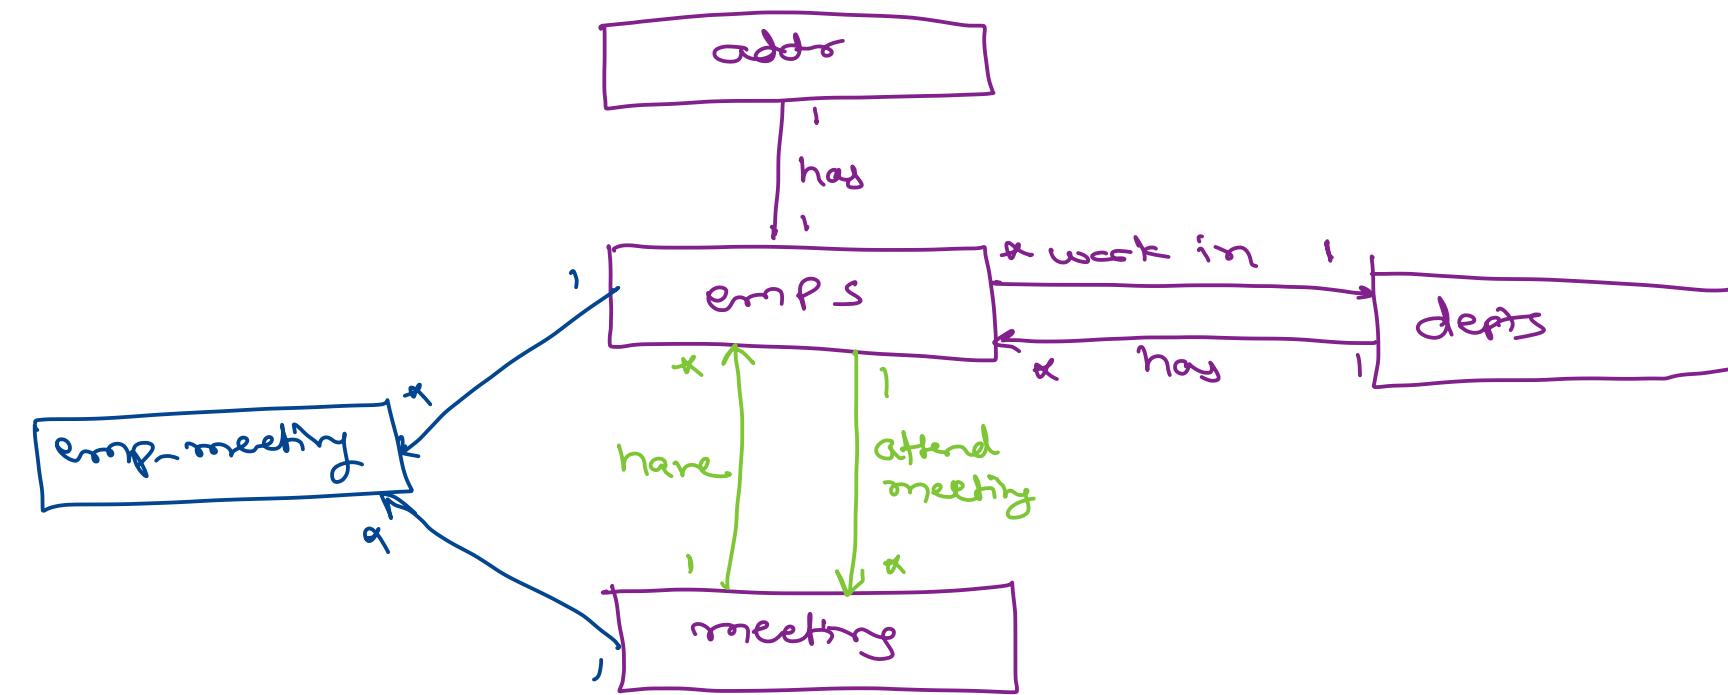
```
select e.ename, d.dname from  
emps e inner join depts d  
on e.deptno = d.deptno;
```

### Non-standard

```
select e.ename, d.dname from  
emps e , depts d  
where e.deptno = d.deptno;
```



# Joins





# Thank you!

Nilesh Ghule <[nilesh@sunbeaminfo.com](mailto:nilesh@sunbeaminfo.com)>





# MySQL - RDBMS

Trainer: Mr. Nilesh Ghule



# Sub queries

- Sub-query is query within query. Typically it work with SELECT statements.
- Output of inner query is used as input to outer query.
- If no optimization is enabled, for each row of outer query result, sub-query is executed once. This reduce performance of sub-query.
- Single row sub-query
  - Sub-query returns single row.
  - Usually it is compared in outer query using relational operators.

where ↗  
having ↗



# Sub queries

- Multi-row sub-query
  - Sub-query returns multiple rows.
  - Usually it is compared in outer query using operators like IN, ANY or ALL.
    - IN operator checks for equality with results from sub-queries. → any of value .
    - ANY operator compares with one of the result from sub-queries. → like OR
    - ALL operator compares with all the results from sub-queries. → like AND





**Thank you!**

Nilesh Ghule <[nilesh@sunbeaminfo.com](mailto:nilesh@sunbeaminfo.com)>





# MySQL - RDBMS

Trainer: Mr. Nilesh Ghule



# Sub queries

- Correlated sub-query

- If number of results from sub-query are reduced, query performance will increase.
- This can be done by adding criteria (WHERE clause) in sub-query based on outer query row.
- Typically correlated sub-query use IN, ALL, ANY and EXISTS operators.

*inner query*



# Sub query

---

- Sub queries with UPDATE and DELETE are not supported in all RDBMS.
  - In MySQL, Sub-queries in UPDATE/DELETE is allowed, but sub-query should not SELECT from the same table, on which UPDATE/DELETE operation is in progress.
- 



# Query performance → RDBMS specific.

- Few RDBMS features ensure better query performance.
  - Index speed up execution of SELECT queries (search operations).
  - Correlated sub-queries execute faster.
- Query performance can be observed using EXPLAIN statement.
  - EXPLAIN FORMAT=JSON SELECT ...;
- EXPLAIN statement shows
  - ✓ Query cost (Lower is the cost, faster is the query execution).
  - ✓ Execution plan (Algorithm used to execute query e.g. loop, semi-join, materialization, etc).
- Optimizations can be enabled or disabled by optimizer\_switch system variable.
  - SELECT @@optimizer\_switch;
  - SET @@optimizer\_switch='materialization=off';

@@  
@  
session variable  
user defined.





**Thank you!**

Nilesh Ghule <[nilesh@sunbeaminfo.com](mailto:nilesh@sunbeaminfo.com)>





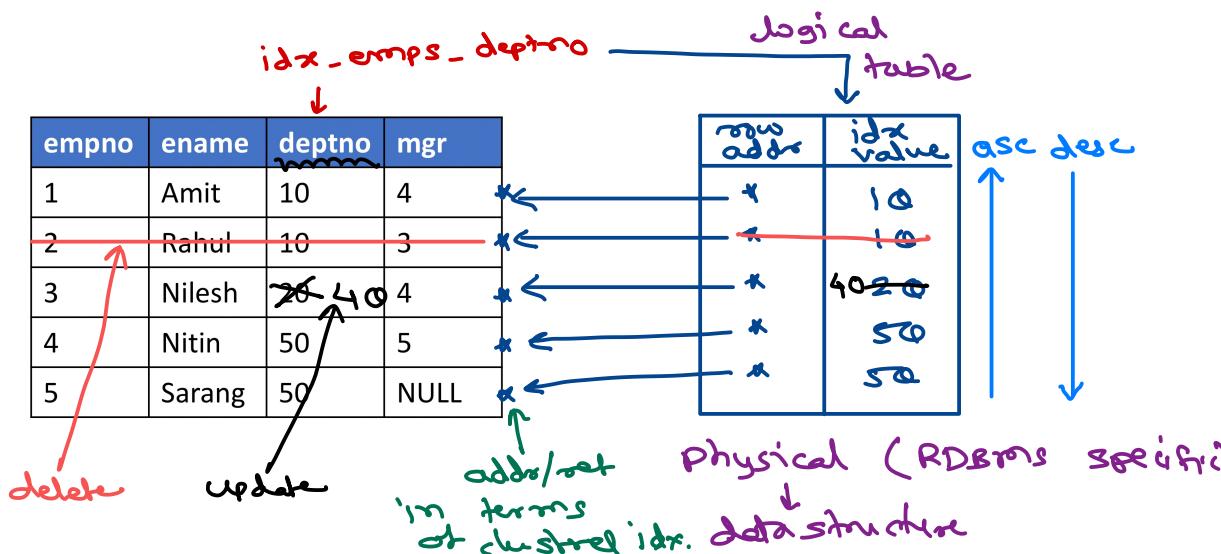
# MySQL - RDBMS

Trainer: Mr. Nilesh Ghule



# Index

- Index enable faster searching in tables by indexed columns.
  - CREATE INDEX idx\_name ON table(column);
- One table can have multiple indexes on different columns/order.
- Typically indexes are stored as some data structure (like BTREE or HASH) on disk.
- Indexes are updated during DML operations. So DML operation are slower on indexed tables.



# Index

- Index can be ASC or DESC.
  - It cause storage of key values in respective order (MySQL 8.x onwards).
  - ASC/DESC index is used by optimizer on ORDER BY queries.
- There are four types of indexes:
  - ✓ Simple index
    - CREATE INDEX idx\_name ON table(column [ASC|DESC]);
  - ✓ Unique index
    - CREATE UNIQUE INDEX idx\_name ON table(column [ASC|DESC]);
    - Doesn't allow duplicate values.
  - ✓ Composite index
    - CREATE INDEX idx\_name ON table(column1 [ASC|DESC], column2 [ASC|DESC]);
    - Composite index can also be unique. Do not allow duplicate combination of columns.
  - ✓ Clustered index
    - PRIMARY index automatically created on Primary key for row lookup.
    - If primary key is not available, hidden index is created on synthetic column. hidden column added by RDBMS
    - It is maintained in tabular form and its reference is used in other indexes.

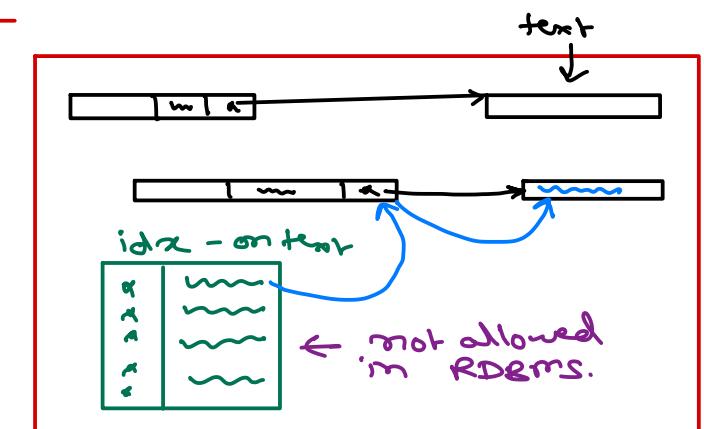


# Index

- Indexes should be created on shorter (INT, CHAR, ...) columns to save disk space.
- Few RDBMS do not allow indexes on external columns i.e. TEXT, BLOB.
- MySQL support indexing on TEXT/BLOB up to n characters.
  - CREATE TABLE test (blob\_col BLOB, ..., INDEX(blob\_col(10)));
- To list all indexes on table:
  - SHOW INDEXES ON table;
- To drop an index:
  - DROP INDEX idx\_name ON table;
- When table is dropped, all indexes are automatically dropped.
- Indexes should not be created on the columns not used frequent search, ordering or grouping operations.
- Columns in join operation should be indexed for better performance.

↓  
↓

index on  
first 6 bytes.  
i.e. only first 10 bytes  
used for searching  
with indexes.



emp(deptno) → ON e.deptno = d.dept no dept (deptno)



# Constraints

- Constraints are **restrictions** imposed on columns.

- There are **five constraints**

- ✓ NOT NULL → col level
- ✓ UNIQUE → col or tbl level
- ✓ PRIMARY KEY → col or tbl level
- ✓ FOREIGN KEY → col or tbl level
- ✓ CHECK → col or tbl level

column value.

create table t1 (

c1 type NOT NULL,  
c2 type unique,  
c3 type,  
c4 type,  
unique(c3),  
constraint consl unique (c3, c4)

);

col level constraint

tbl level constraint

auto generated name for constraint

- Few constraints can be applied at either **column level** or **table level**. Few constraints can be applied on both.
- Optionally constraint names can be mentioned while creating the constraint. If not given, it is auto-generated.
- Each DML operation check the constraints before manipulating the values. If any constraint is violated, error is raised.



# Constraints

- **NOT NULL**

- NULL values are not allowed.
- Can be applied at column level only.
- CREATE TABLE table(c1 TYPE NOT NULL, ...);

- **UNIQUE**

- Duplicate values are not allowed.
- NULL values are allowed.
- Not applicable for TEXT and BLOB.
- UNIQUE can be applied on one or more columns.
- Internally creates unique index on the column (fast searching).
- Can be applied at column level or table level.
  - CREATE TABLE table(c1 TYPE UNIQUE, ...);
  - CREATE TABLE table(c1 TYPE, ..., UNIQUE(c1));
  - CREATE TABLE table(c1 TYPE, ..., CONSTRAINT constraint\_name UNIQUE(c1));

```
Create table students (
    std INT not null,
    roll INT not null,
    name CHAR(20),
    unique (std, roll)
);
```

internally create unique  
Composite index.





**Thank you!**

Nilesh Ghule <[nilesh@sunbeaminfo.com](mailto:nilesh@sunbeaminfo.com)>





# MySQL - RDBMS

Trainer: Mr. Nilesh Ghule



# Constraints

- NOT NULL

- NULL values are not allowed.
- Can be applied at column level only.
- CREATE TABLE table(c1 TYPE NOT NULL, ...);

- UNIQUE

- Duplicate values are not allowed.
- NULL values are allowed.
- Not applicable for TEXT and BLOB.
- UNIQUE can be applied on one or more columns.
- Internally creates unique index on the column (fast searching).
- Can be applied at column level or table level.
  - CREATE TABLE table(c1 TYPE UNIQUE, ...);
  - CREATE TABLE table(c1 TYPE, ..., UNIQUE(c1));
  - CREATE TABLE table(c1 TYPE, ..., CONSTRAINT constraint\_name UNIQUE(c1));

one table may have multiple unique keys.

```
Create table students (
    std INT not null,
    roll INT not null,
    name CHAR(20),
    unique (std, roll)
);
```

internally create unique  
Composite index.



# Constraints

*internally used for clustered index.*

- **PRIMARY KEY**

- Column or set of columns that uniquely identifies a row.
- Only one primary key is allowed for a table.
- Primary key column cannot have duplicate or NULL values. = unique + not null.
- Internally index is created on PK column.
- TEXT/BLOB cannot be primary key.
- If no obvious choice available for PK, composite or surrogate PK can be created.
- Creating PK for a table is a good practice.
- PK can be created at table level or column level.
- CREATE TABLE table(c1 TYPE PRIMARY KEY, ...);
- CREATE TABLE table(c1 TYPE, ..., PRIMARY KEY(c1));
- CREATE TABLE table(c1 TYPE, ..., CONSTRAINT constraint\_name PRIMARY KEY(c1));
- CREATE TABLE table(c1 TYPE, c2 TYPE, ..., PRIMARY KEY(c1, c2));

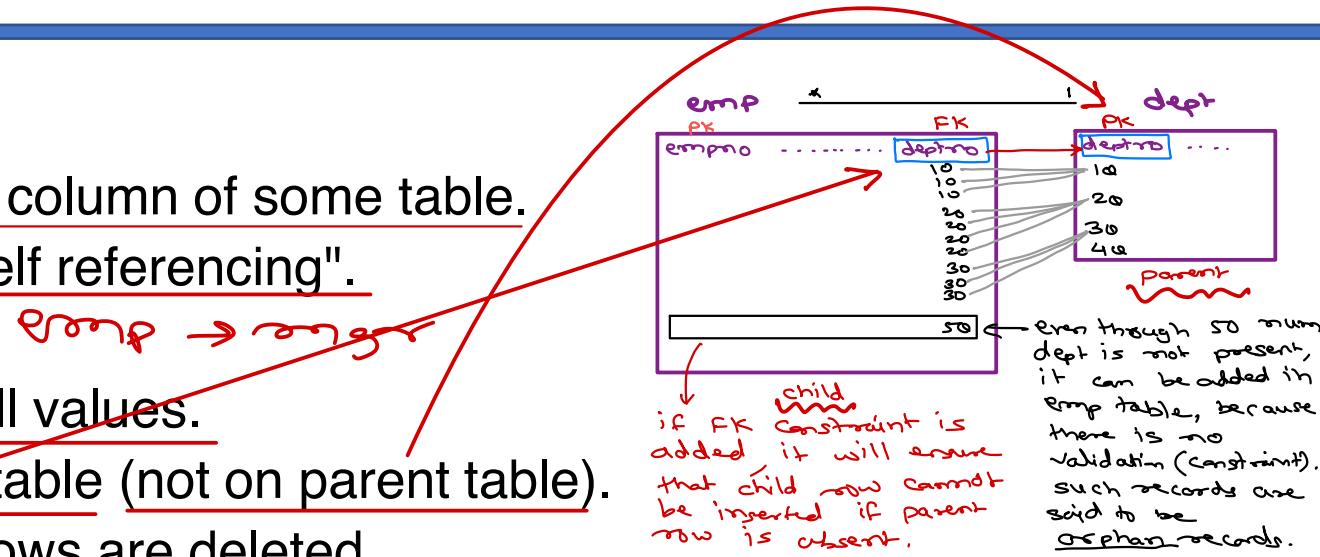


# Constraints

## • FOREIGN KEY

*composite FK*

- Column or set of columns that references a column of some table.
- If column belongs to the same table, it is "self referencing".
- FK can have duplicate values as well as null values.
- FK constraint is applied on column of child table (not on parent table).
- parent rows cannot be deleted, until child rows are deleted.
- MySQL have ON DELETE CASCADE clause to ensure that child rows are automatically deleted, when parent row is deleted. ON UPDATE CASCADE clause does same for UPDATE operation.
- By default foreign key checks are enabled. They can be disabled by
  - `SET @@foreign_key_checks = 0;`
  - FK constraint can be applied on table level as well as column level.
  - `CREATE TABLE child(c1 TYPE, ..., FOREIGN KEY (c1) REFERENCES parent(col))`



# Constraints

- **CHECK**
  - CHECK is integrity constraint in SQL.
  - CHECK constraint specifies condition on column.
  - Data can be inserted/updated only if condition is true; otherwise error is raised.
  - CHECK constraint can be applied at table level or column level.
  - CREATE TABLE table(c1 TYPE, c2 TYPE CHECK condition1, ..., CHECK condition2);





**Thank you!**

Nilesh Ghule <[nilesh@sunbeaminfo.com](mailto:nilesh@sunbeaminfo.com)>





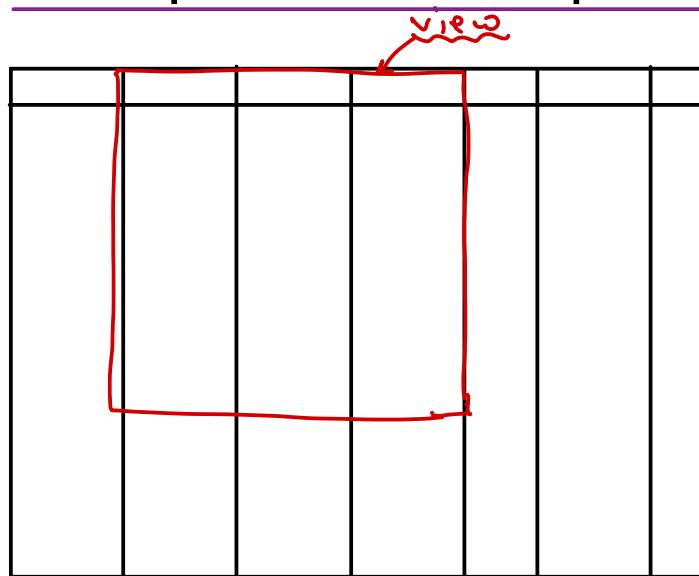
# MySQL - RDBMS

Trainer: Mr. Nilesh Ghule



# Views

- RDBMS view represents view (projection) of the data.
- View is based on SELECT statement.
- Typically it is restricted view of the data (limited rows or columns) from one or more tables (joins and/or sub-queries) or summary of the data (grouping).
- Data of view is not stored on server hard-disk; but its SELECT statement is stored in compiled form. It speed up execution of view.



→ each query on view will internally execute select query first. & process on top of it.

create view view-name  
AS SELECT ----- ;

oracle                    views                    mysql  
materialized            non-materialized  
view                      view  
view data is stored    view data is  
in temp memory.        not stored.

# Views

Cannot perform

DML operations

- Views are of two types: Simple view and Complex view
- Usually if view contains computed columns, group by, joins or sub-queries, then the views are said to be complex. DML operations are not supported on these views.
- DML operations on view affects underlying table.
- View can be created with CHECK OPTION to ensure that DML operations can be performed only the data visible in that view.



# View

- Views can be differentiated with: SHOW FULL TABLES.
- Views can be dropped with DROP VIEW statement.
- View can be based on another view.

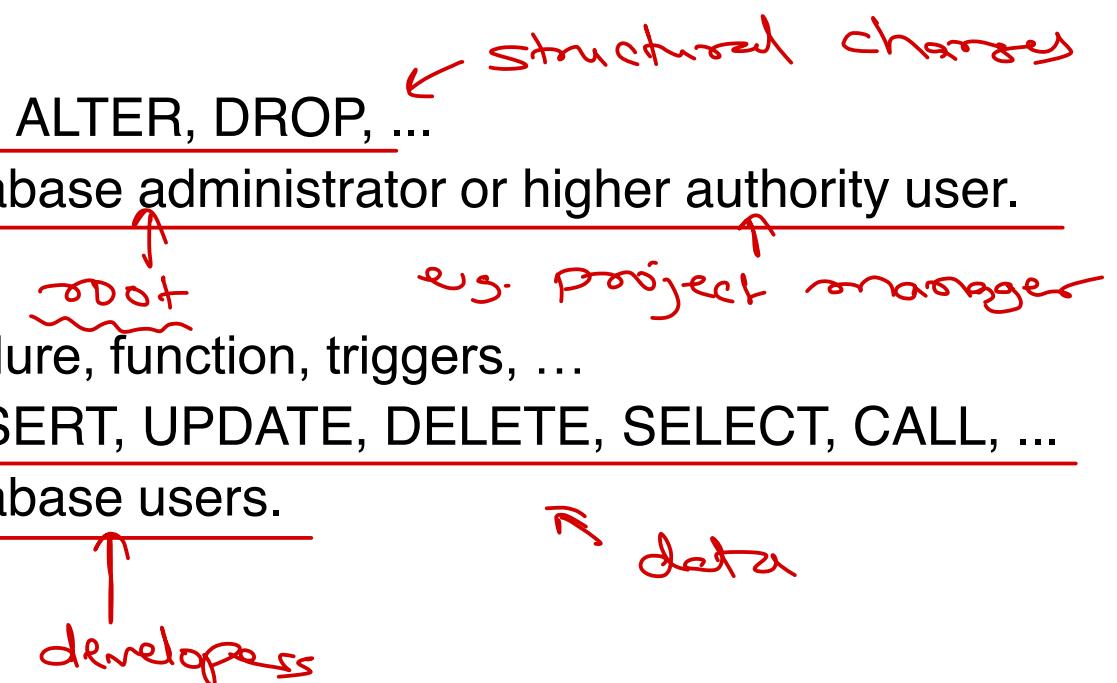
*create view view2 as  
Select \* from view1 where —;*

- Applications of views
  - ✓ Security: Providing limited access to the data.
  - ✓ Hide source code of the table.
  - ✓ Simplifies complex queries.



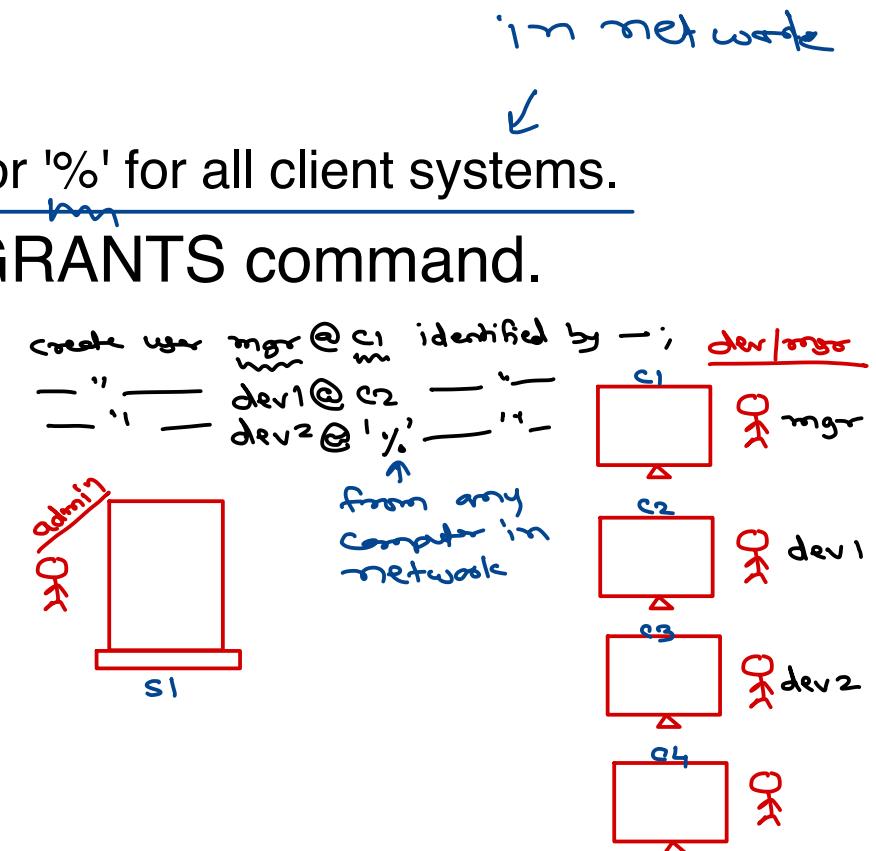
# Data Control Language

- Security is built-in feature of any RDBMS. It is implemented in terms of permissions (a.k.a. privileges).
- There are two types of privileges.
- System privileges
  - Privileges for certain commands i.e. CREATE, ALTER, DROP, ...
  - Typically these privileges are given to the database administrator or higher authority user.
- Object privileges
  - RDBMS objects are table, view, stored procedure, function, triggers, ...
  - Can perform operations on the objects i.e. INSERT, UPDATE, DELETE, SELECT, CALL, ...
  - Typically these privileges are given to the database users.



# User Management

- User management is responsibility of admin (root).
  - New user can be created using CREATE USER.
    - CREATE USER user@host IDENTIFIED BY 'password';
    - host can be hostname of server, localhost (current system) or '%' for all client systems.
  - Permissions for the user can be listed using SHOW GRANTS command.
    - SHOW GRANTS FOR user@host;
  - Users can be deleted using DROP USER.
    - DROP USER user@host;
  - Change user password.
    - ALTER USER user@host IDENTIFIED BY 'new\_password';
    - FLUSH PRIVILEGES;



# Data Control Language

- Permissions are given to user using GRANT command.
  - GRANT CREATE ON db.\* TO user@host; *global - on all db.*
  - GRANT CREATE ON \*.\* TO user1@host, user2@host;
  - GRANT SELECT ON db.table TO user@host;
  - GRANT SELECT, INSERT, UPDATE ON db.table TO user@host;
  - GRANT ALL ON db.\* TO user@host;
- By default one user cannot give permissions to other user. This can be enabled using WITH GRANT OPTION. *similar to "root".*
  - GRANT ALL ON \*.\* TO user@host WITH GRANT OPTION;
- Permissions assigned to any user can be withdrawn using REVOKE command.
  - REVOKE SELECT, INSERT ON db.table FROM user@host;
- Permissions can be activated by FLUSH PRIVILEGES.
  - System GRANT tables are reloaded by this command. Auto done after GRANT, REVOKE.
  - Command is necessary if GRANT tables are modified using DML operations.



# DDL – ALTER statement

- ALTER statement is used to do modification into table, view, function, procedure, ...
- ALTER TABLE is used to change table structure.
- Add new column(s) into the table.
  - ALTER TABLE table ADD col TYPE;
  - ALTER TABLE table ADD c1 TYPE, c2 TYPE;
- Modify column of the table.
  - ALTER TABLE table MODIFY col NEW\_TYPE;
- Rename column.
  - ALTER TABLE CHANGE old\_col new\_col TYPE;
- Drop a column
  - ALTER TABLE DROP COLUMN col;
- Rename table
  - ALTER TABLE table RENAME TO new\_table;



# PSM - Agenda - Persistent Storage Module -

- MySQL Programming
- ✓ • Stored procedure
- ✓ • Exceptions
- ✓ • Function
- ✓ • Trigger

Program → set of instructions

MySQL program → set of SQL statements.  
along with programming  
constructs e.g.  
loops, if-else, --.



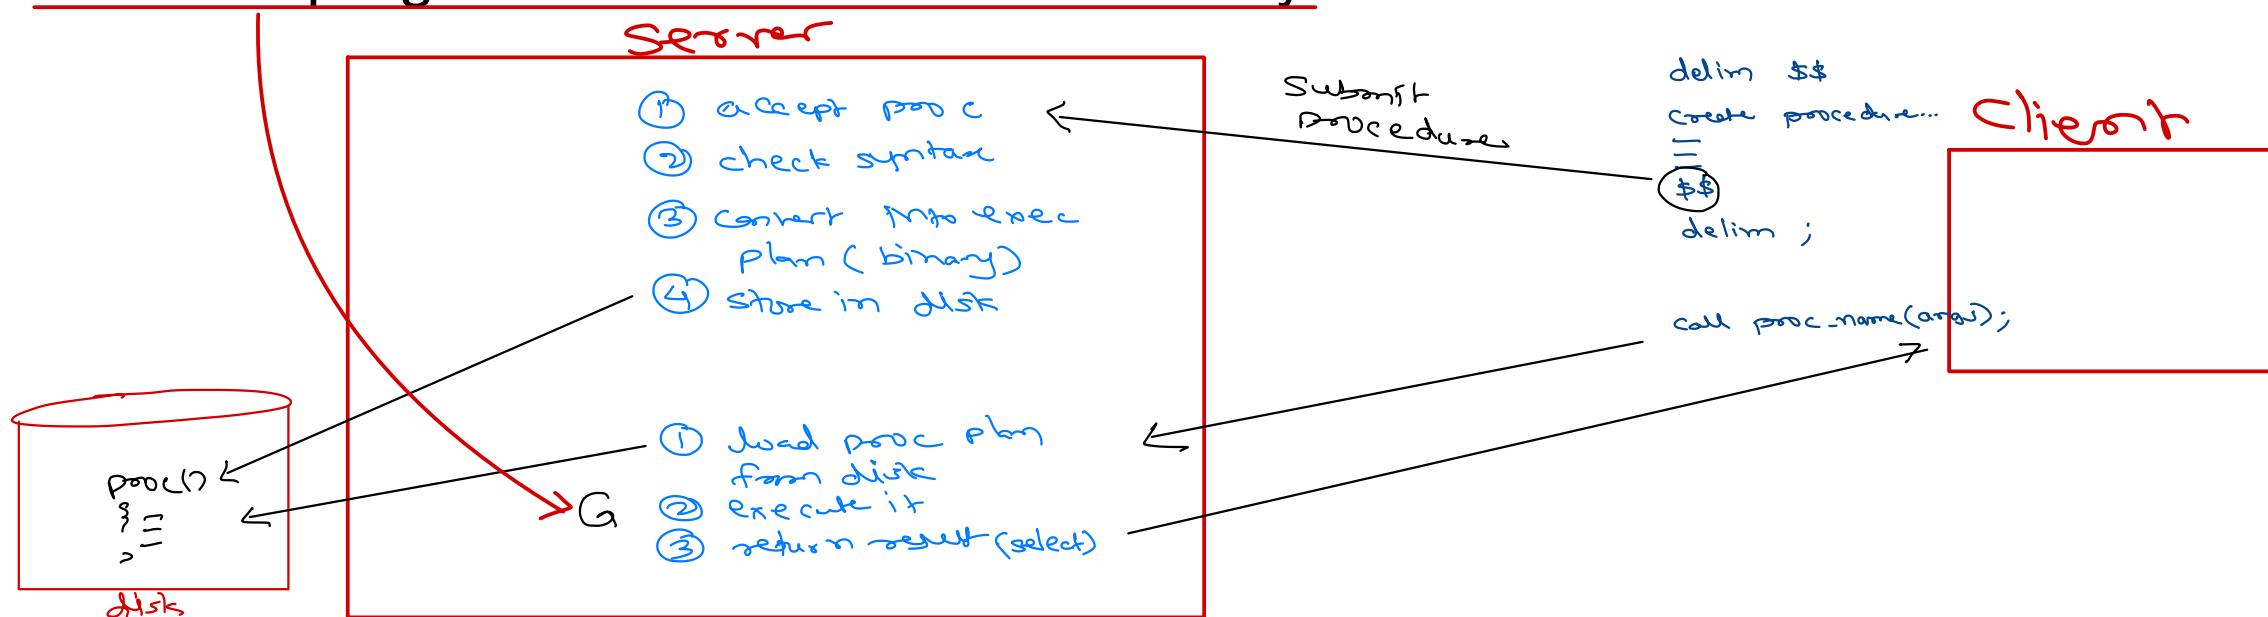
# MySQL Programming

- RDBMS Programming is an ISO standard – part of SQL standard – since 1992.
- SQL/PSM stands for Persistent Stored Module.
- Inspired from PL/SQL - Programming language of Oracle.
- PSM allows writing programs for RDBMS. The program contains set of SQL statements along with programming constructs e.g. variables, if-else, loops, case, ...
- PSM is a block language. Blocks can be nested into another block.
- MySQL program can be a stored procedure, function or trigger.



# MySQL Programming

- MySQL PSM program is written by db user (programmers).
- It is submitted from client, server check syntax & store them into db in compiled form.
- The program can be executed by db user when needed.
- Since programs are stored on server in compiled form, their execution is very fast.
- All these programs will run in server memory.



# Stored Procedure

- Stored Procedure is a routine. It contains multiple SQL statements along with programming constructs.
- Procedure doesn't return any value (like void fns in C).
- Procedures can take zero or more parameters.
- Procedures are created using CREATE PROCEDURE and deleted using DROP PROCEDURE.
- Procedures are invoked/called using CALL statement.
- Result of stored procedure can be
  - returned via OUT parameter.
  - inserted into another table.
  - produced using SELECT statement (at end of SP).
- Delimiter should be set before writing SQL query.



# Stored Procedure

```
CREATE TABLE result(v1 DOUBLE, v2 VARCHAR(50));          -- 01_hello.sql (using editor)
DELIMITER $$                                           DROP PROCEDURE IF EXISTS sp_hello;
                                                       DELIMITER $$

CREATE PROCEDURE sp_hello()                           CREATE PROCEDURE sp_hello()
BEGIN                                              BEGIN
    INSERT INTO result VALUES(1, 'Hello World');   SELECT 1 AS v1, 'Hello World' AS v2;
END;                                                 END;

$$                                                 $$

DELIMITER ;                                         DELIMITER ;
                                                       SOURCE /path/to/01_hello.sql
CALL sp_hello();                                     CALL sp_hello();

SELECT * FROM result;
```



# Stored Procedure – PSM Syntax

## VARIABLES

```
DECLARE varname DATATYPE;  
DECLARE varname DATATYPE DEFAULT init_value;  
SET varname = new_value;  
SELECT new_value INTO varname;  
SELECT expr_or_col INTO varname FROM table_name;
```

## PARAMETERS

```
CREATE PROCEDURE sp_name(PARAMTYPE p1 DATATYPE)  
BEGIN  
...  
END;  
  
-- IN param: Initialized by calling program.  
-- OUT param: Initialized by called procedure.  
-- INOUT param: Initialized by calling program and  
modified by called procedure  
-- OUT & INOUT param declared as session variables.  
  
CREATE PROCEDURE sp_name(OUT p1 INT)  
BEGIN  
    SELECT 1 INTO p1;  
END;  
  
SET @res = 0;  
CALL sp_name(@res);  
SELECT @res;
```

## IF-ELSE

```
IF condition THEN  
    body;  
END IF;  
-----  
IF condition THEN  
    if-body;  
ELSE  
    else-body;  
END IF;  
-----  
IF condition THEN  
    if1-body;  
ELSE  
    IF condition THEN  
        if2-body;  
    ELSE  
        else2-body;  
    END IF;  
END IF;  
-----  
IF condition THEN  
    if1-body;  
ELSEIF condition THEN  
    if2-body;  
ELSE  
    else-body;  
END IF;
```

## LOOPS

```
WHILE condition DO  
    body;  
END WHILE;  
-----  
REPEAT  
    body;  
UNTIL condition  
END REPEAT;  
-----  
label: LOOP  
IF condition THEN  
    ...  
    LEAVE label;  
END IF;  
...  
END LOOP;
```

## SHOW PROCEDURE

```
SHOW PROCEDURE STATUS  
LIKE 'sp_name';  
  
SHOW CREATE PROCEDURE sp_name;
```

## DROP PROCEDURE

```
DROP PROCEDURE  
IF EXISTS sp_name;
```

## CASE-WHEN

```
CASE  
WHEN condition THEN  
    body;  
WHEN condition THEN  
    body;  
ELSE  
    body;  
END CASE;
```





**Thank you!**

Nilesh Ghule <[nilesh@sunbeaminfo.com](mailto:nilesh@sunbeaminfo.com)>





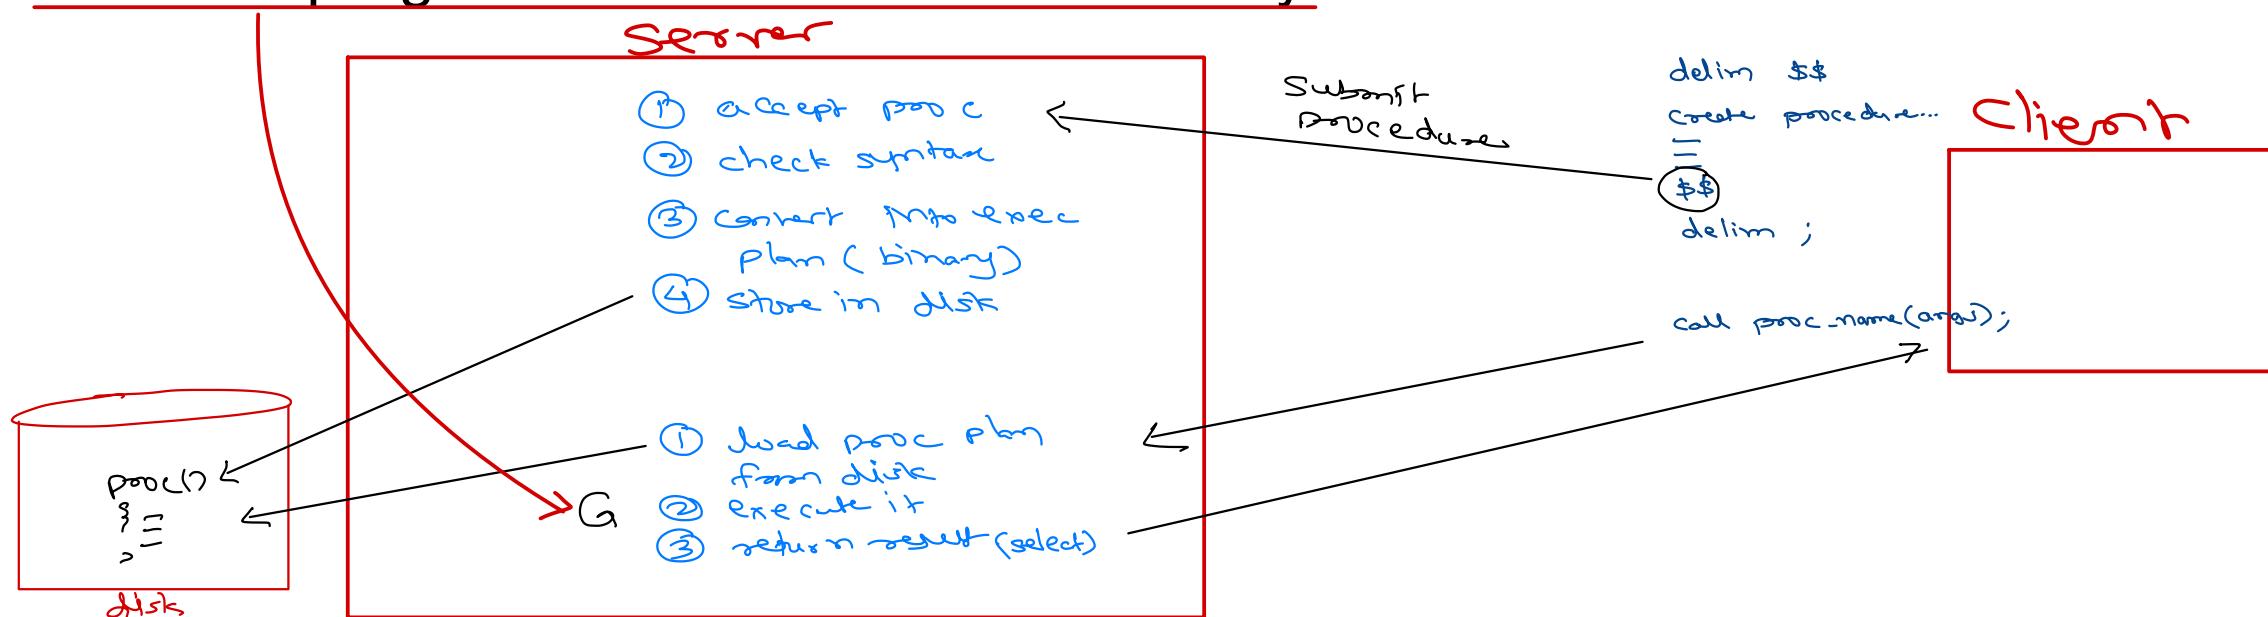
# MySQL - RDBMS

Trainer: Mr. Nilesh Ghule



# MySQL Programming

- MySQL PSM program is written by db user (programmers).
- It is submitted from client, server check syntax & store them into db in compiled form.
- The program can be executed by db user when needed.
- Since programs are stored on server in compiled form, their execution is very fast.
- All these programs will run in server memory.



# Stored Procedure

- Stored Procedure is a routine. It contains multiple SQL statements along with programming constructs.
- Procedure doesn't return any value (like void fns in C).
- Procedures can take zero or more parameters.
- Procedures are created using CREATE PROCEDURE and deleted using DROP PROCEDURE.
- Procedures are invoked/called using CALL statement.
- Result of stored procedure can be
  - returned via OUT parameter.
  - inserted into another table. → e.g. results table
  - produced using SELECT statement (at end of SP) → only last select statement will be displayed.
- Delimiter should be set before writing SQL query.

procedure



# Stored Procedure

```
CREATE TABLE result(v1 DOUBLE, v2 VARCHAR(50));  
DELIMITER $$  
  
CREATE PROCEDURE sp_hello()  
BEGIN  
    INSERT INTO result VALUES(1, 'Hello World');  
END;  
$$  
  
DELIMITER ;  
  
CALL sp_hello();  
?  
? SELECT * FROM result; ← see the result.
```

```
-- 01_hello.sql (using editor)  
DROP PROCEDURE IF EXISTS sp_hello;  
DELIMITER $$  
CREATE PROCEDURE sp_hello()  
BEGIN  
    ✓ SELECT 1 AS v1, 'Hello World' AS v2;  
END;  
$$  
DELIMITER ;
```

SOURCE /path/to/01\_hello.sql

CALL sp\_hello();

forward slash  
no space in whole path.  
Output will be displayed here.(CLI).



# Stored Procedure – PSM Syntax

## VARIABLES

```
DECLARE varname DATATYPE;  
DECLARE varname DATATYPE DEFAULT init_value;  
SET varname = new_value;  
SELECT new_value INTO varname;  
SELECT expr_or_col INTO varname FROM table_name;
```

## PARAMETERS

```
CREATE PROCEDURE sp_name(PARAMTYPE p1 DATATYPE)  
BEGIN  
...  
END;  
  
-- IN param: Initialized by calling program.  
-- OUT param: Initialized by called procedure.  
-- INOUT param: Initialized by calling program and  
modified by called procedure  
-- OUT & INOUT param declared as session variables.  
  
CREATE PROCEDURE sp_name(OUT p1 INT)  
BEGIN  
    SELECT 1 INTO p1;  
END;  
  
SET @res = 0;  
CALL sp_name(@res);  
SELECT @res;
```

## IF-ELSE

```
IF condition THEN  
    body;  
END IF;  
-----  
IF condition THEN  
    if-body;  
ELSE  
    else-body;  
END IF;  
-----  
IF condition THEN  
    if1-body;  
ELSE  
    IF condition THEN  
        if2-body;  
    ELSE  
        else2-body;  
    END IF;  
END IF;  
-----  
IF condition THEN  
    if1-body;  
ELSEIF condition THEN  
    if2-body;  
ELSE  
    else-body;  
END IF;
```

## LOOPS

```
WHILE condition DO  
    body;  
END WHILE;  
-----  
REPEAT  
    body;  
UNTIL condition  
END REPEAT;  
-----  
label: LOOP  
IF condition THEN  
    ...  
    LEAVE label;  
END IF;  
...  
END LOOP;
```

## SHOW PROCEDURE

```
SHOW PROCEDURE STATUS  
LIKE 'sp_name';  
  
SHOW CREATE PROCEDURE sp_name;
```

## DROP PROCEDURE

```
DROP PROCEDURE  
IF EXISTS sp_name;
```

## CASE-WHEN

```
CASE  
WHEN condition THEN  
    body;  
WHEN condition THEN  
    body;  
ELSE  
    body;  
END CASE;
```



# MySQL Exceptions / Error Handling

- Exceptions are runtime problems, which may arise during execution of stored procedure, function or trigger.
- Required actions should be taken against these errors.
- SP execution may be continued or stopped after handling exception.
- MySQL error handlers are declared as:
  - DECLARE action HANDLER FOR condition handler\_impl;
  - The *action* can be: CONTINUE or EXIT.
  - The *condition* can be:
    - MySQL error code: e.g. 1062 for duplicate entry.
    - SQLSTATE value: e.g. 23000 for duplicate entry, NOTFOUND for end-of-cursor.
    - Named condition: e.g. DECLARE duplicate\_entry CONDITION FOR 1062;
  - The *handler\_impl* can be: Single liner or PSM block i.e. BEGIN ... END;

```
create procedure sp_div(v_num int, v_den int)
begin
    declare v_res int default 0;
    set v_res = v_num / v_den; ? 
    select v_res as result;
```

error/exception  
end;





**Thank you!**

Nilesh Ghule <[nilesh@sunbeaminfo.com](mailto:nilesh@sunbeaminfo.com)>





# MySQL - RDBMS

Trainer: Mr. Nilesh Ghule



# MySQL Stored Functions

information\_schema → routines table  
SP  
FN

- Stored Functions are MySQL programs like stored procedures.
- Functions can be having ~~zero~~ or more parameters. MySQL allows only ~~IN~~ params.
- Functions must return some value using RETURN statement.
- Function entire code is stored in system table.
- Like procedures, functions allows statements like local variable declarations, if-else, case, loops, etc. One function can invoke another function/procedure and vice-versa.  
The functions can also be recursive.
- There are two types of functions: DETERMINISTIC and NOT DETERMINISTIC.

## CREATE FUNCTION

```
CREATE FUNCTION fn_name(p1 TYPE)
RETURNS TYPE
[NOT] DETERMINISTIC
BEGIN
    body;
    RETURN value;
END;
```

## SHOW FUNCTION

```
SHOW FUNCTION STATUS LIKE 'fn_name';
SHOW CREATE FUNCTION fn_name;
```

## DROP FUNCTION

```
DROP FUNCTION IF EXISTS fn_name;
```

for given param values, return value always remain same.

even if input param are same, return value may differ.  
in this case result also depend on external factor like date time, state table, ...



# MySQL Triggers → also psm program —Same syntax

- Triggers are supported by all standard RDBMS like Oracle, MySQL, etc.
- Triggers are not supported by WEAK RDBMS like MS-- Access. SQLite, ...
- Triggers are not called by client's directly, so they don't have args & return value.
- Trigger execution is caused by DML operations on database.
  - BEFORE/AFTER INSERT, BEFORE/AFTER UPDATE, BEFORE/AFTER DELETE.
- Like SP/FN, Triggers may contain SQL statements with programming constructs. They may also call other SP or FN.
- However COMMIT/ROLLBACK is not allowed in triggers. They are executed in same transaction in which DML query is executed.

## CREATE TRIGGER

```
CREATE TRIGGER trig_name
AFTER|BEFORE dml_op ON table
FOR EACH ROW
BEGIN
    body;
    -- use OLD & NEW keywords
    -- to access old/new rows.
    -- INSERT triggers - NEW rows.
    -- DELETE triggers - OLD rows.
END;
```

## SHOW TRIGGERS

```
SHOW TRIGGERS FROM db_name;
```

## DROP TRIGGER

```
DROP TRIGGER trig_name;
```

SP → CALL sp-name();

FN → SELECT fn-name()  
...;

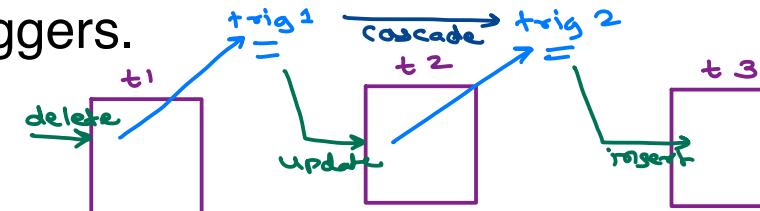
Trigger → Not call.



# MySQL Triggers

- Applications of triggers:
  - ✓ Maintain logs of DML operations (Audit Trails).
  - ✓ Data cleansing before insert or update data into table. (Modify NEW value).
  - ✓ Copying each record AFTER INSERT into another table to maintain "Shadow table".
  - ✓ Copying each record AFTER DELETE into another table to maintain "History table".
  - ✓ Auto operations of related tables using cascading triggers.

↑ backup

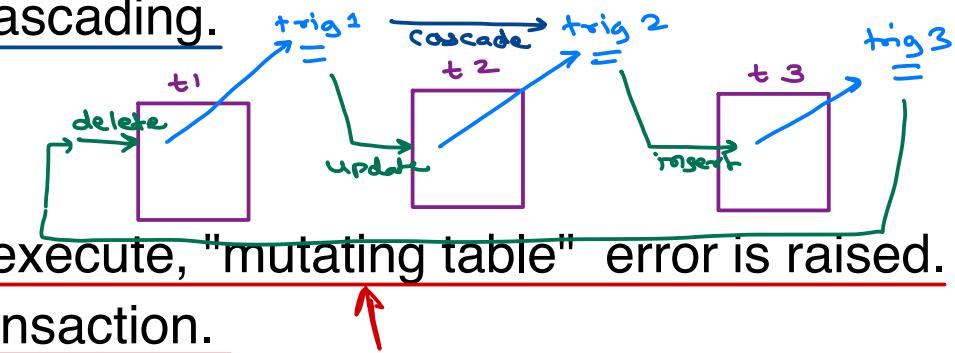


## Cascading triggers

- One trigger causes execution of 2<sup>nd</sup> trigger, 2<sup>nd</sup> trigger causes execution of 3<sup>rd</sup> trigger and so on.
- In MySQL, there is no upper limit on number of levels of cascading.
- This is helpful in complicated business processes.

## Mutating table error

- If cascading trigger causes one of the earlier trigger to re-execute, "mutating table" error is raised.
- This prevents infinite loop and also rollback the current transaction.





**Thank you!**

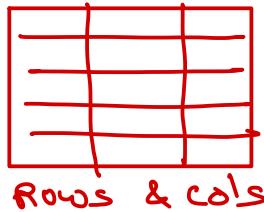
Nilesh Ghule <[nilesh@sunbeaminfo.com](mailto:nilesh@sunbeaminfo.com)>



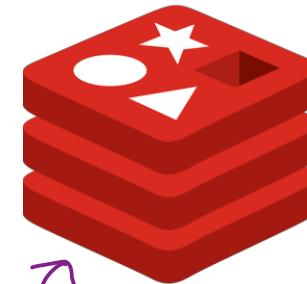
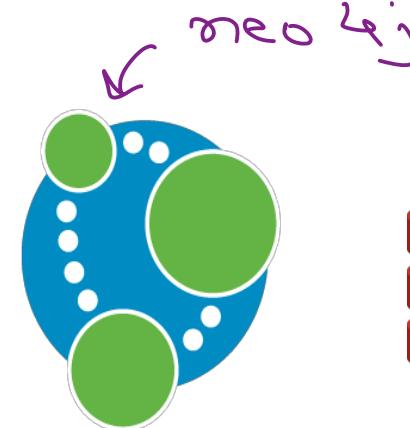
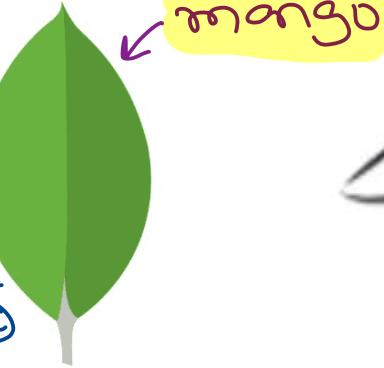
all RD**BMS** → SQL

Relational

tabular [ $T_1 \leftrightarrow T_2$ ]



- ✓ fixed structure (schema)
- ✓ 100s of GB
- ✓ high consistency / transactions.



- ✓ flexible schema
- ✓ 100s of TB/PB (scaling)
- ✓ data
- ✓ economical
- ✓ eventual consistency

## NoSQL Databases

Trainer: Mr. Nilesh Ghule

Not Only SQL

→ no standard query language X

# Document oriented databases

Java Script Object Notation

- Document contains data as key-value pair as **JSON** or XML.
- Document schema is flexible & are added in collection for processing.
- RDBMS tables → Collections
- RDBMS rows → Documents
- RDBMS columns → Key-value pairs in document
- Examples: MongoDb, CouchDb, ...

b1 JSON → Java Script Object Notation.  
{  
  "id": 1, → int  
  "title": "Let us C", → string  
  "author": "Karnetkar",  
  "price": 240.4 → double  
}  
3

r1 → JSON document  
{  
  id: 1,  
  name: "Nilesh",  
  age: 38,  
  hobbies: ["Program", "Reading", ...]  
  addr: { area: "Katraj", city: "Pune", pin: 411046 },  
  Political: false,  
  height: 5.9,  
  bloodgroup: null  
}





**mongoDB**<sup>®</sup>

## MongoDb Databases

Trainer: Mr. Nilesh Ghule



Sunbeam Infotech

[www.sunbeaminfo.com](http://www.sunbeaminfo.com)

# Mongo Db

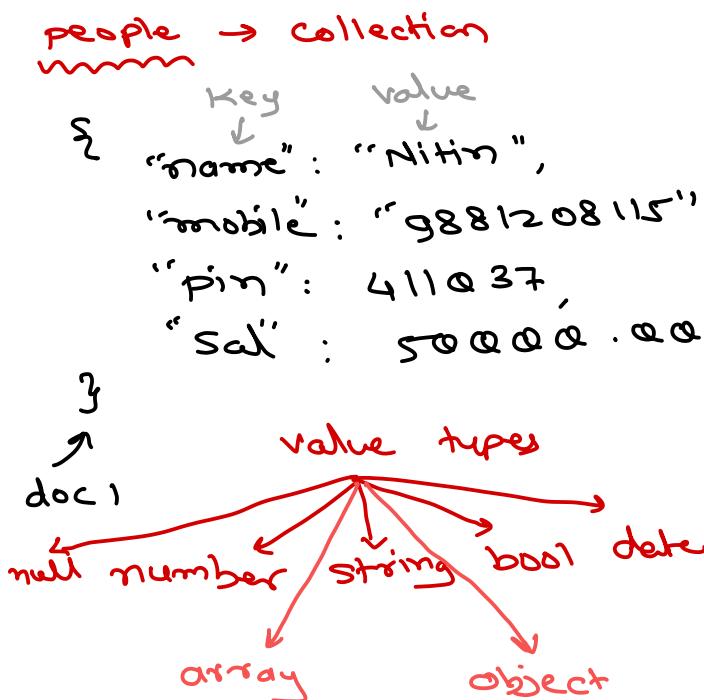
- Developed by 10gen in 2007
- Publicly available in 2009
- Open-source database which is controlled by 10gen
- Document oriented database → stores JSON documents
- Stores data in binary JSON. → (BSON) → faster access.
- Design Philosophy
  - ✓ MongoDB wasn't designed in a lab and is instead built from the experiences of building large scale, high availability, and robust systems.  
PBs      24x7



# JSON

- Java Script Object Notation
- Hierarchical way of organizing data
- Mongo stores JSON data into Binary form.

→ data type internally identified as a number.



doc 2 →

{

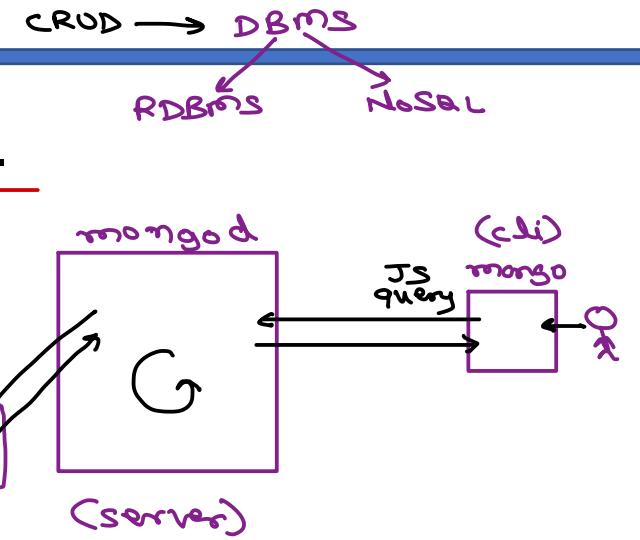
    "name": "Nilesh", ← string  
    "add": { ← json object  
        "area": "katraj",  
        "city": "pune",  
        "pin": 411046  
    },  
    "age": 37, ← array  
    "sal": 300000.0,  
    "political": false, ← boolean  
    "rating": null ← null

}



# Mongo Server and Client

- MongoDb server (mongod) is developed in C, C++ and JS.
- MongoDb data is accessed via multiple client tools
  - ✓ mongo : client shell (JS).
  - ✓ mongofiles : stores larger files in GridFS.
  - ✓ mongoimport / mongoexport : tools for data import / export.
  - ✓ mongodump / mongorestore : tools for backup / restore.
- MongoDb data can be accessed in application through client drivers available for all major programming languages e.g. Java, Python, Ruby, PHP, Perl, ...
- Mongo shell is follows JS syntax and allow to execute JS scripts.  
~~~~~



# MongoDb: Data Types

| data    | bson        | values                                         |
|---------|-------------|------------------------------------------------|
| null    | 10          |                                                |
| boolean | 8           | true, false                                    |
| number  | 1 / 16 / 18 | 123, 456.78, NumberInt("24"), NumberLong("28") |
| string  | 2           | "...."                                         |
| date    | 9           | new Date(), ISODate("yyyy-mm-ddThh:mm:ss")     |
| array   | 4           | [ ..., ..., ..., ... ]                         |
| object  | 3           | { ... }                                        |



# Mongo - INSERT

- show databases;
- use database;
- db.contacts.insert({name: "nilesh", mobile: "9527331338"});
- db.contacts.insertMany([  
    {name: "nilesh", mobile: "9527331338"},  
    {name: "nitin", mobile: "9881208115"}  
]);
- Maximum document size is 16 MB.
- For each object unique id is generated by client (if \_id not provided).
  - 12 byte unique id :: [counter(3) I pid(2) I machine(3) I timestamp(4)]



# Mongo – QUERY

- db.contacts.find(); → returns cursor on which following ops allowed:
  - hasNext(), next(), skip(n), limit(n), count(), toArray(), forEach(fn), pretty()
- Shell restrict to fetch 20 records at once. Press "it" for more records.
- db.contacts.find( { name: "nilesh" } );
- db.contacts.find( { name: "nilesh" }, { \_id:0, name:1 } );
- Relational operators: \$eq, \$ne, \$gt, \$lt, \$gte, \$lte, \$in, \$nin
- Logical operators: \$and, \$or, \$nor, \$not
- Element operators: \$exists, \$type
- Evaluation operators: \$regex, \$where, \$mod
- Array operators: \$size, \$elemMatch, \$all, \$slice





# Thank you!

Nilesh Ghule <[nilesh@sunbeaminfo.com](mailto:nilesh@sunbeaminfo.com)>

